

Robotics 2

TORO - Efficient Constraint Network Optimization for SLAM

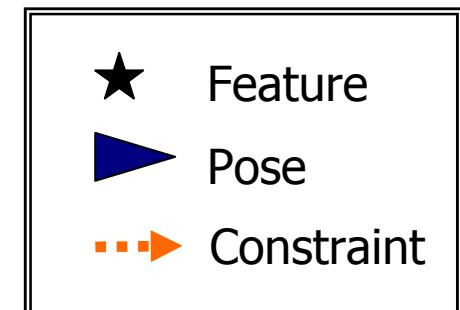
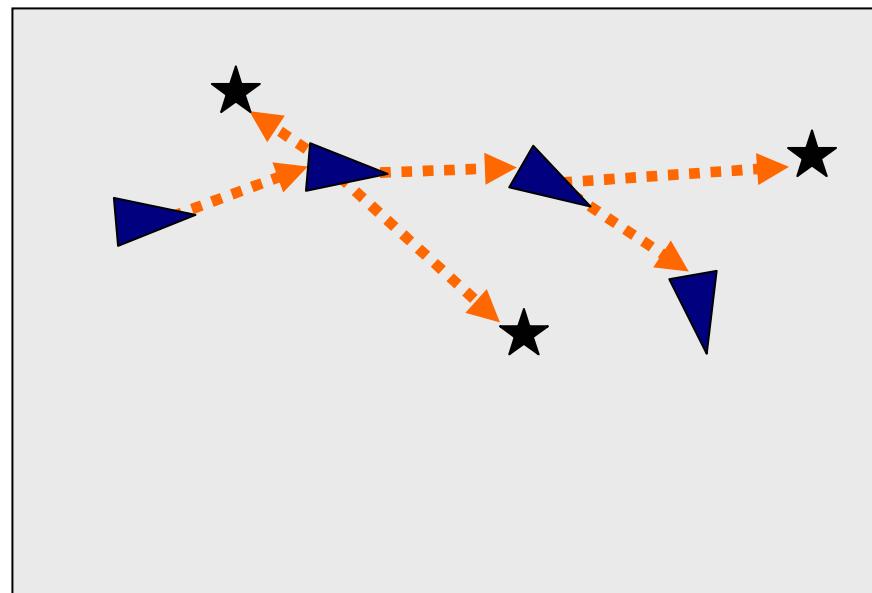
Giorgio Grisetti, Cyrill Stachniss,

Kai Arras, Wolfram Burgard



Robot Mapping

- Constraints connect the poses of the robot while it is moving (odometry).

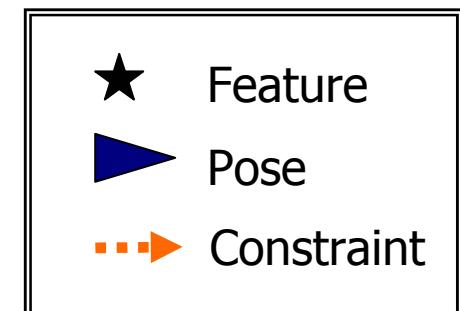
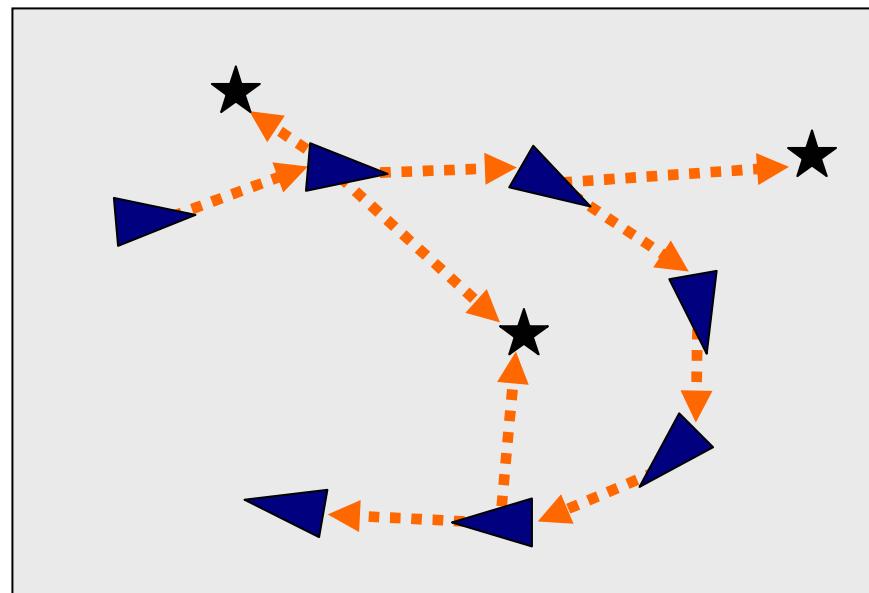


State: x, y, yaw

or $x, y, z, \text{roll}, \text{pitch}, \text{yaw}$

Robot Mapping

- Re-observing features defines constraints between non-successive poses.



State: x, y, yaw

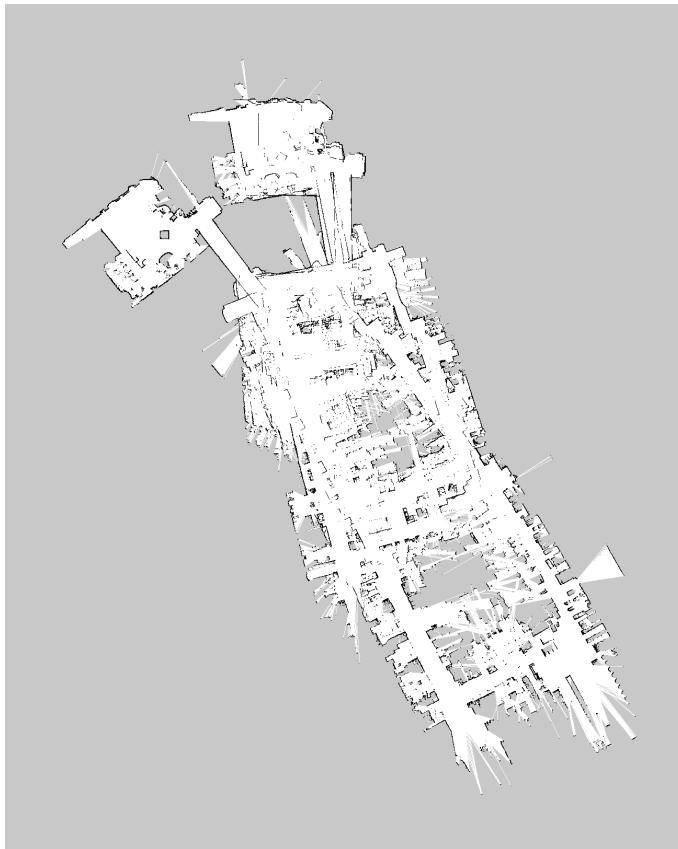
or $x, y, z, \text{roll}, \text{pitch}, \text{yaw}$

Graph-based SLAM

- SLAM = simultaneous localization and mapping
- Use a graph to represent the problem
- Every node in the graph corresponds to a pose of the robot during mapping
- Every edge between two nodes corresponds to the spatial constraints between them
- **Goal:** Find a configuration of the nodes that minimize the error introduced by the constraints

Problem Summary

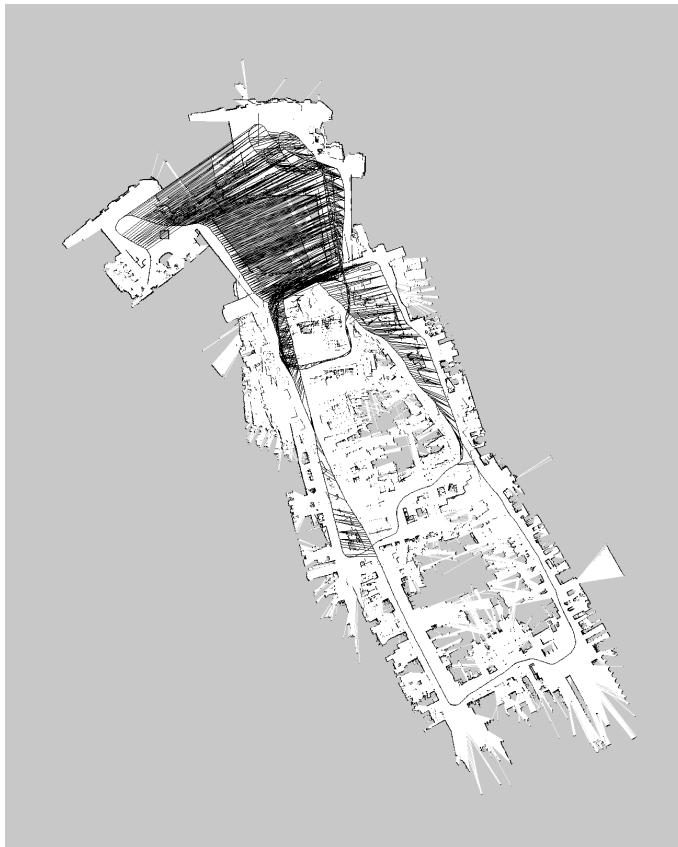
- **Goal:** Find the arrangement of the nodes that satisfies the constraints best



An initial configuration (KUKA production hall 22)

Problem Summary

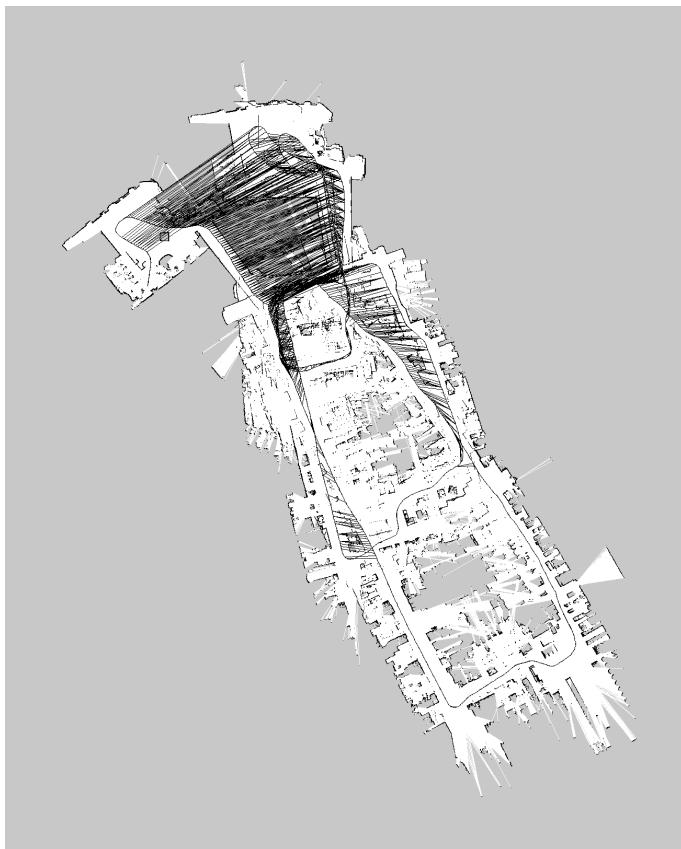
- **Goal:** Find the arrangement of the nodes that satisfies the constraints best



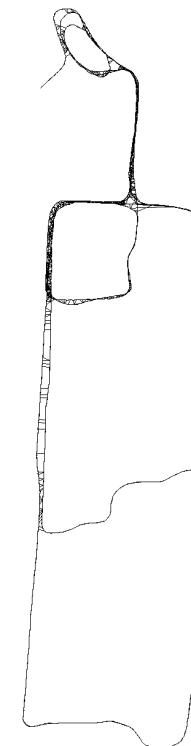
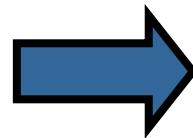
An initial configuration (KUKA production hall 22)

Problem Summary

- **Goal:** Find the arrangement of the nodes that satisfies the constraints best



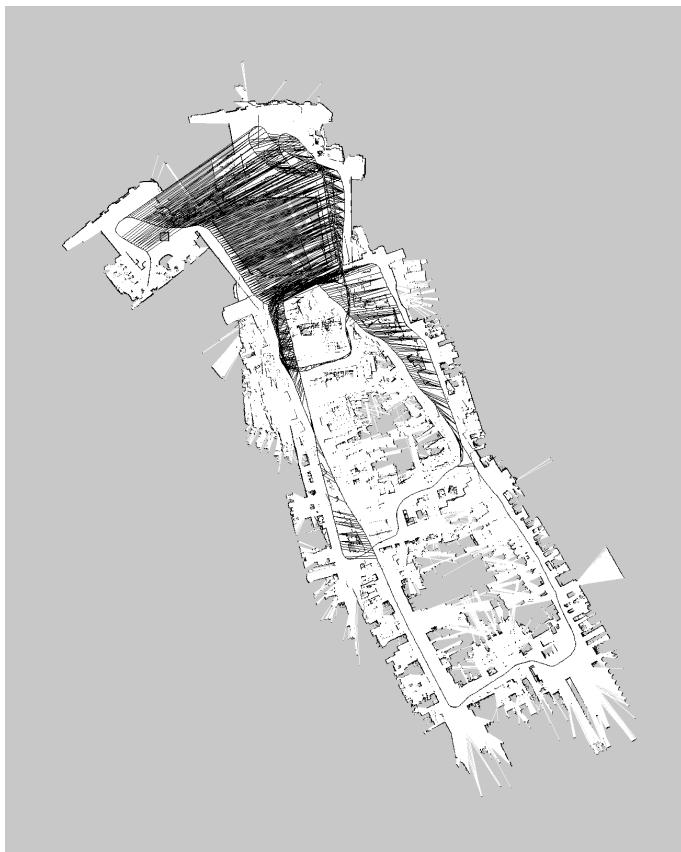
An initial configuration



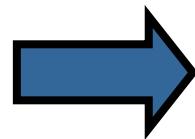
Maximum likelihood configuration

Problem Summary

- **Goal:** Find the arrangement of the nodes that satisfies the constraints best

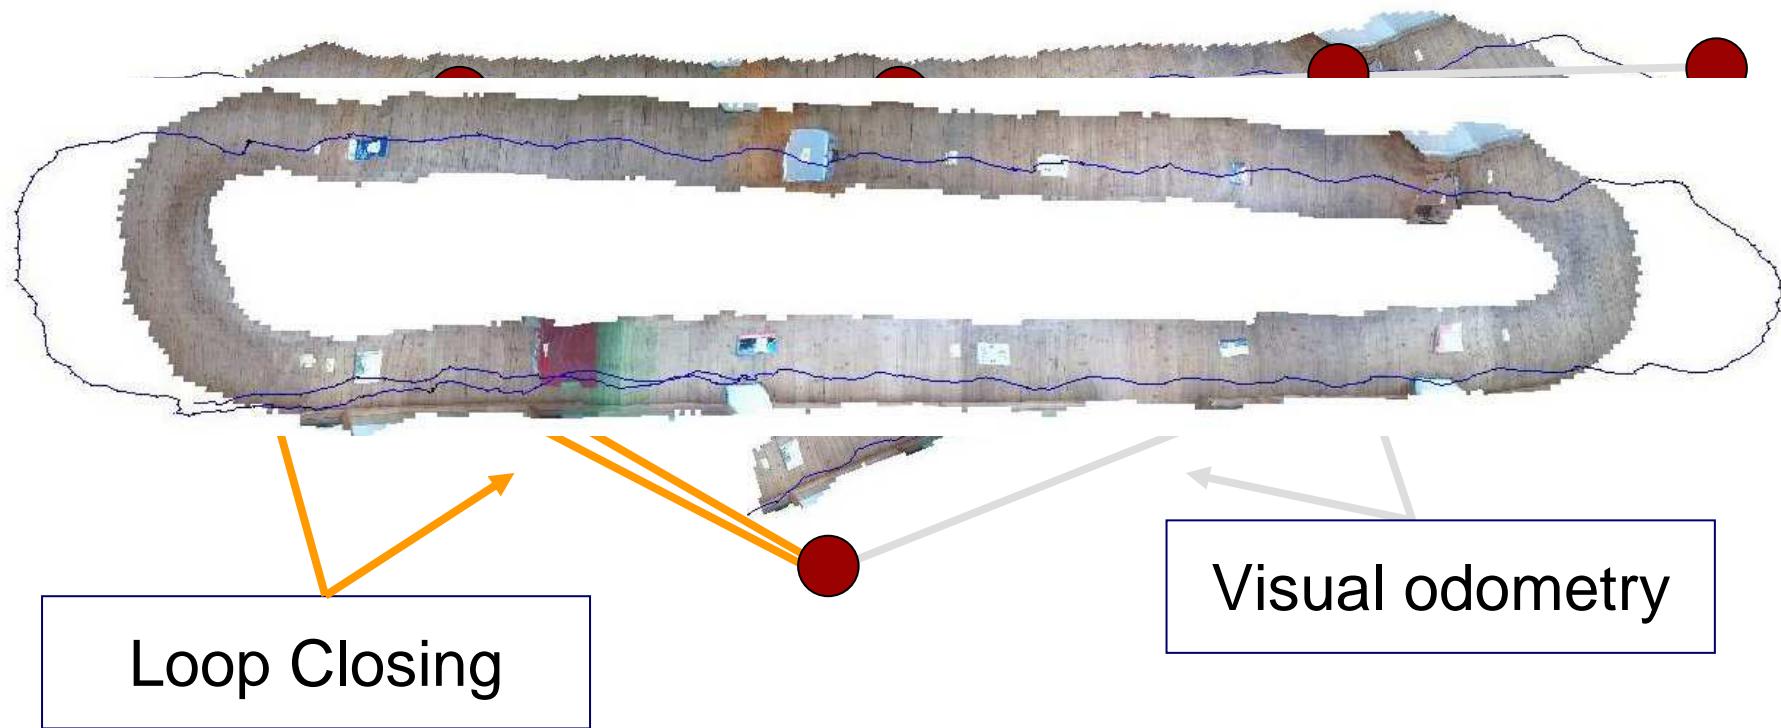


An initial configuration



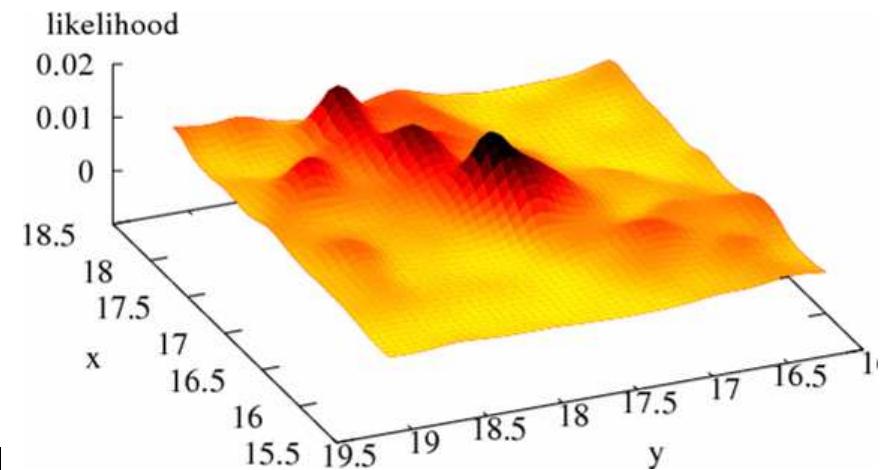
Maximum likelihood configuration

Graph-based Visual SLAM



Why is This Hard?

- Huge number of unknowns
 - $|x| = 10^3$ is a *small* problem
- Non-linear constraints
 - Function of the orientation of the robot
- Bad initial estimates
 - Odometry (if at all)
- Local minima
 - in the observation likelihood $p(z | x)$
- Possibility to incremental computation the solution



Topics Today

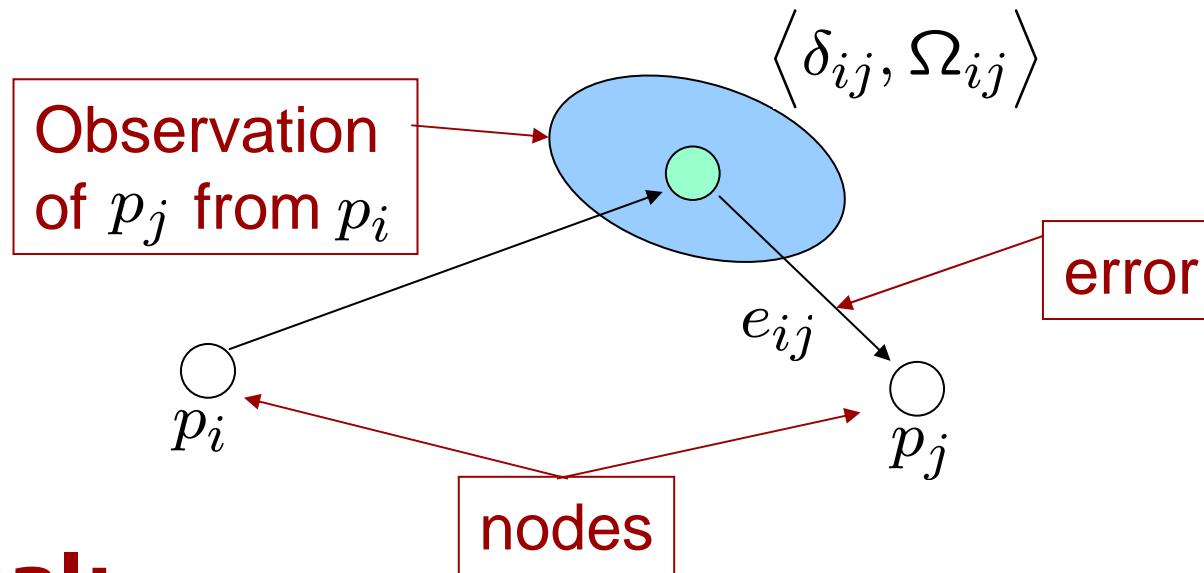
- Estimate the Gaussian posterior about the poses of the robot (full SLAM)

Three Parts:

- Estimate the means via non-linear optimization (maximum likelihood map)
- Estimate the covariance matrices via belief propagation and covariance intersection
- Examples of how to obtain constraints

Problem Formulation

- The problem can be described by a graph



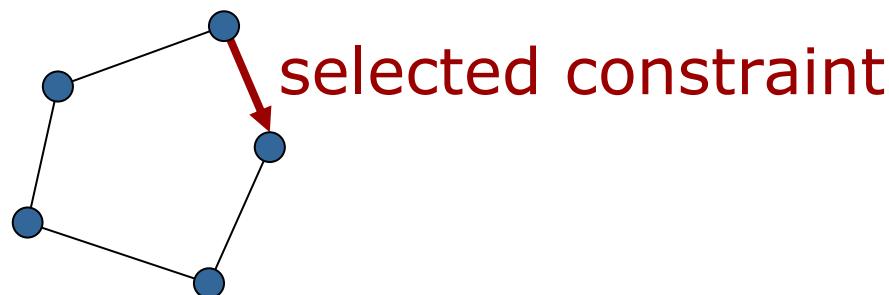
Goal:

- Find the assignment of poses to the nodes of the graph which minimizes the negative log likelihood of the observations:

$$\hat{\mathbf{p}} = \operatorname{argmin} \sum_{ij} e_{ij}^T \Omega_{ij} e_{ij}$$

Stochastic Gradient Descent

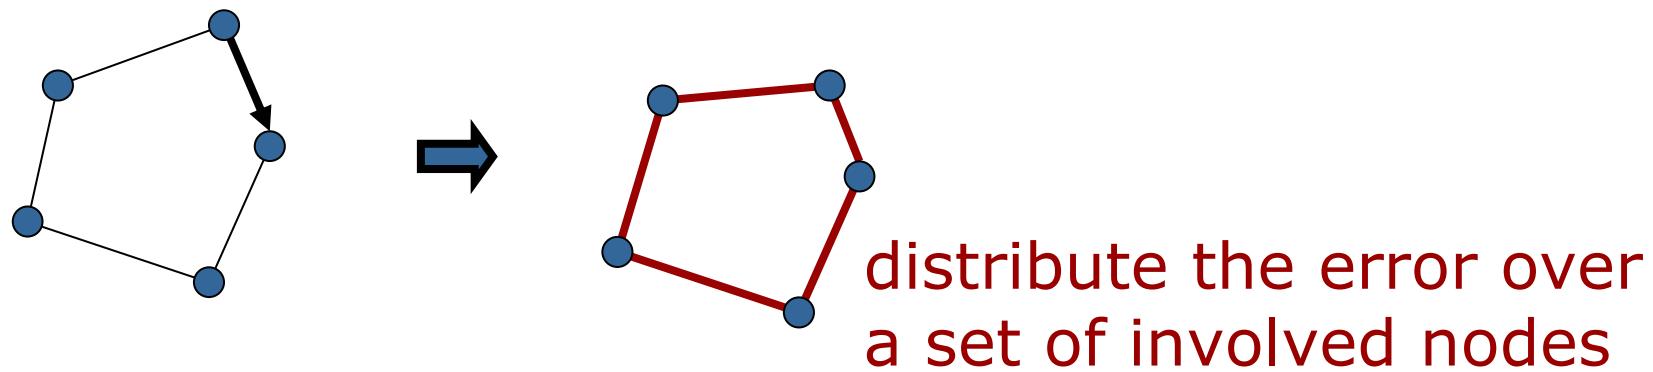
- Minimize the error individually for each constraint (decomposition of the problem into sub-problems)
- Solve one step of each sub-problem
- Solutions might be contradictory
- The magnitude of the correction decreases with each iteration
- Learning rate to achieve convergence



[First introduced in the SLAM community by Olson et al., '06]

Stochastic Gradient Descent

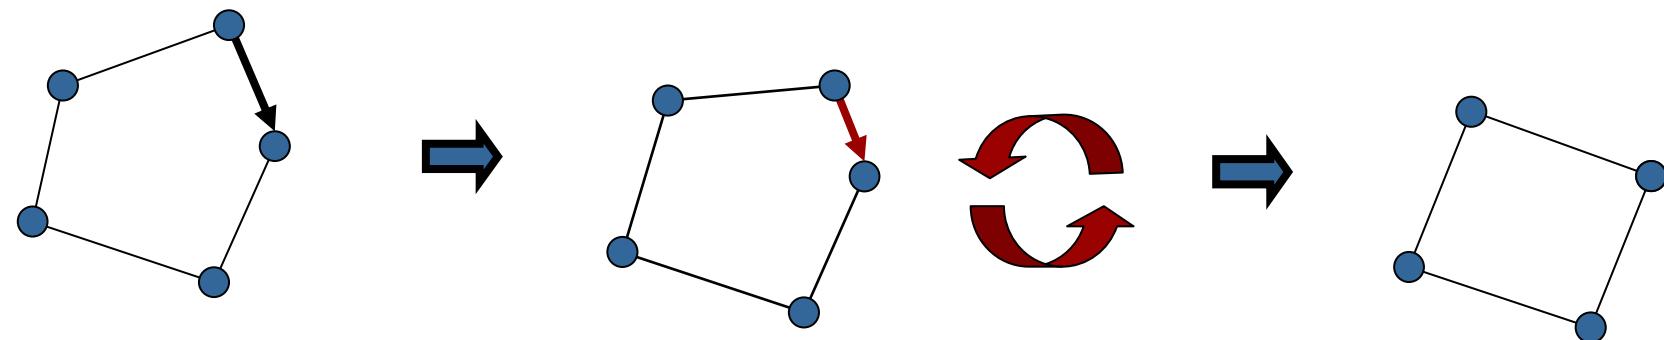
- Minimize the error individually for each constraint (decomposition of the problem into sub-problems)
- Solve one step of each sub-problem
- Solutions might be contradictory
- The magnitude of the correction decreases with each iteration
- Learning rate to achieve convergence



[First introduced in the SLAM community by Olson et al., '06]

Stochastic Gradient Descent

- Minimize the error individually for each constraint (decomposition of the problem into sub-problems)
- Solve one step of each sub-problem
- Solutions might be contradictory
- The magnitude of the correction decreases with each iteration
- Learning rate to achieve convergence



[First introduced in the SLAM community by Olson et al., '06]

Preconditioned SGD

- Minimize the error individually for each constraint (decomposition of the problem into sub-problems)
- Solve one step of each sub-problem
- Solutions might be contradictory
- A solution is found when an equilibrium is reached
- Update rule for a single constraint:

$$\mathbf{x}^{t+1} = \mathbf{x}^t + \lambda \cdot \mathbf{H}^{-1} J_{ij}^T \Omega_{ij} r_{ij}$$

Diagram illustrating the components of the update rule:

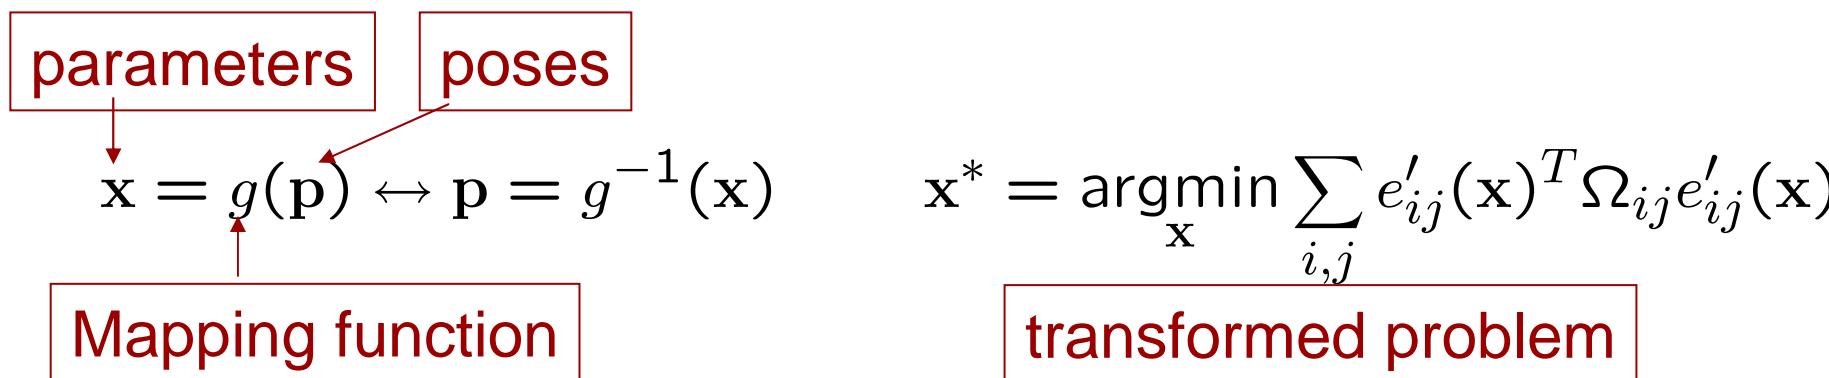
Previous solution	Hessian	Information matrix
\mathbf{x}^t	\mathbf{H}^{-1}	$J_{ij}^T \Omega_{ij} r_{ij}$
Current solution	Learning rate	Jacobian residual

The diagram shows the flow of information from the "Current solution" and "Learning rate" (both pointing to \mathbf{x}^t) through the "Jacobian" and "residual" (both pointing to $J_{ij}^T \Omega_{ij} r_{ij}$) to the final result \mathbf{x}^{t+1} .

[First introduced in the SLAM community by Olson et al., '06]

Node Parameterization

- How to represent the nodes in the graph?
- **Key question: which parts need to be updated for a single constraint update?**
- This are to the “sub-problems” in SGD
- Transform the problem into a different space so that:
 - the structure of the problem is exploited.
 - the calculations become easier and faster.



Parameterization of Olson

- Incremental parameterization:

$$x_i = p_i - p_{i-1}$$

The diagram illustrates the incremental parameterization formula $x_i = p_i - p_{i-1}$. It features two red-bordered boxes: one labeled "parameters" and another labeled "poses". Red arrows point from these boxes to the terms p_i and p_{i-1} in the equation, respectively.

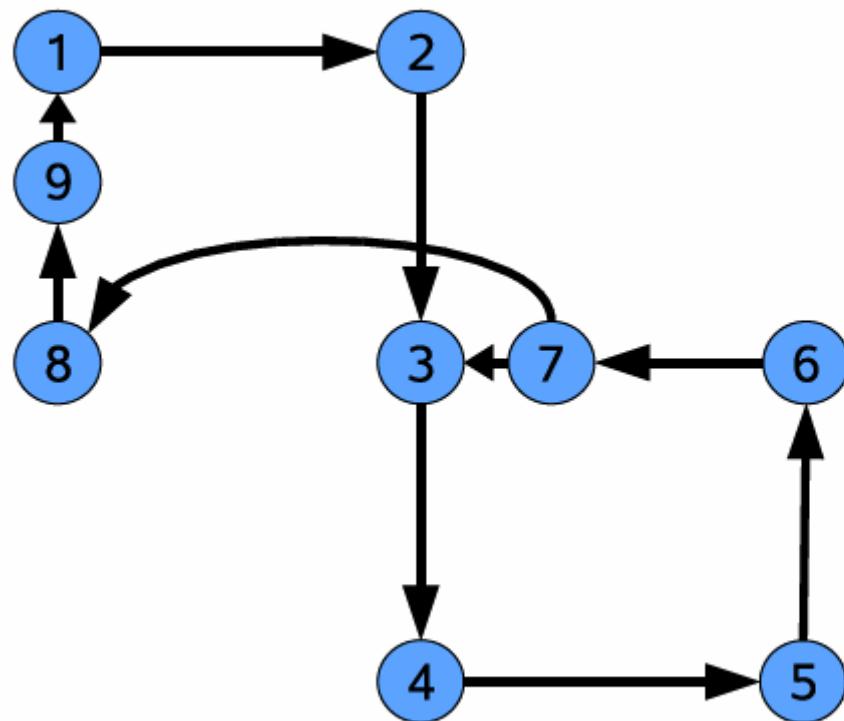
- Results directly from the trajectory takes by the robot
- Problem: for optimizing a constraint between node i and k , one needs to updates all $j = i, \dots, k$ nodes ignoring the topology of the environment

Alternative Parameterization

- Exploit the topology of the space to compute the parameterization
- Idea: “Loops should be one sub-problem”
- Such a parameterization can be extracted from the graph topology itself

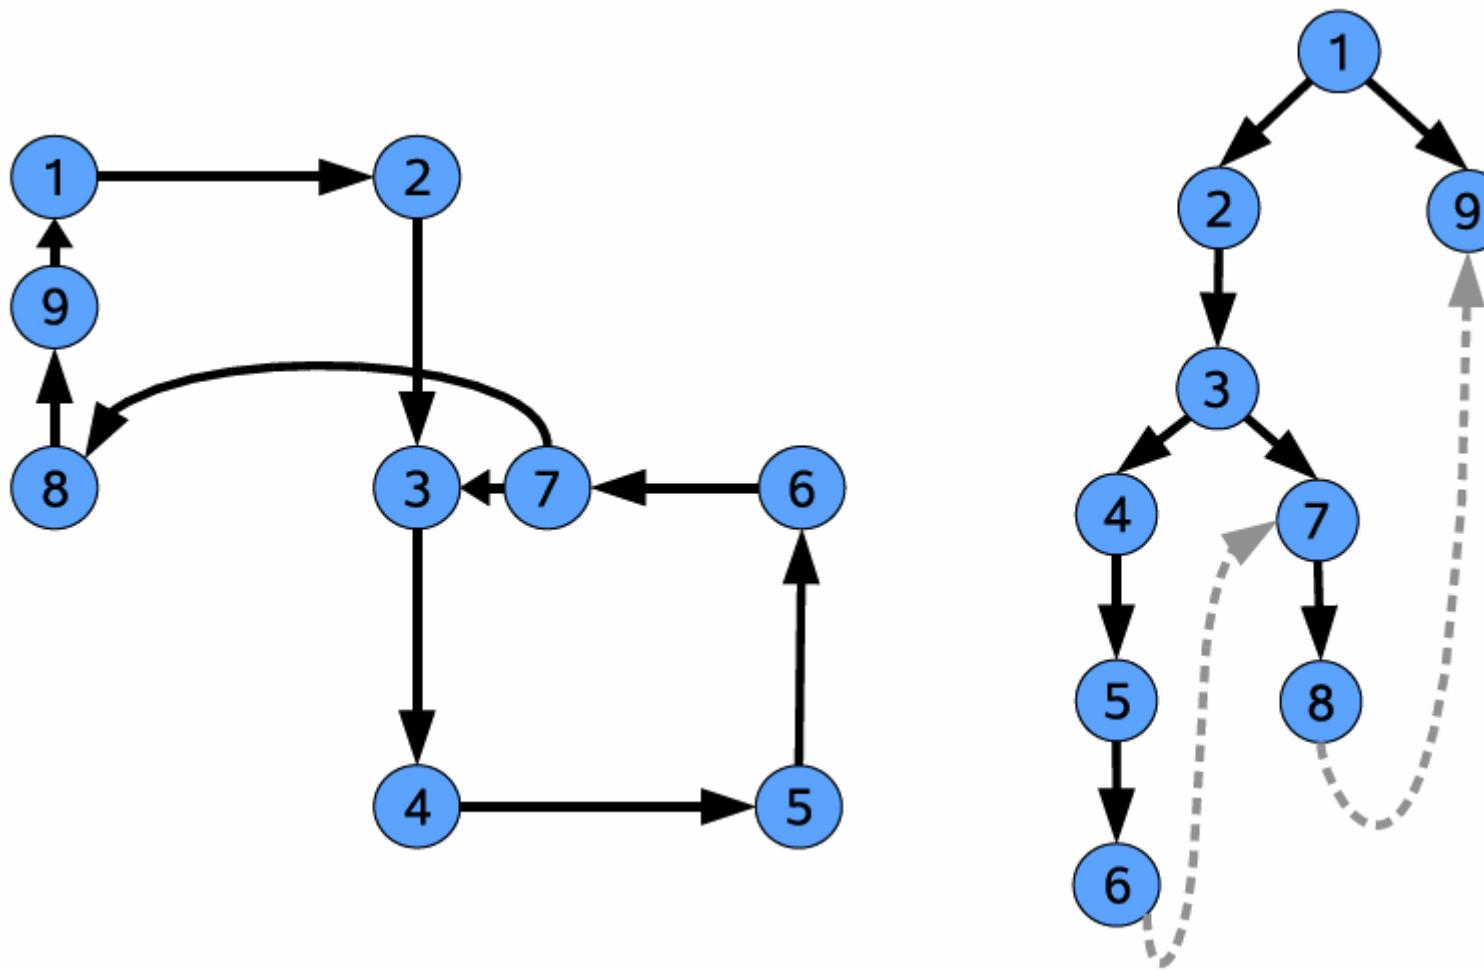
Tree Parameterization

- How should such a problem decomposition look like?



Tree Parameterization

- Use a spanning tree!

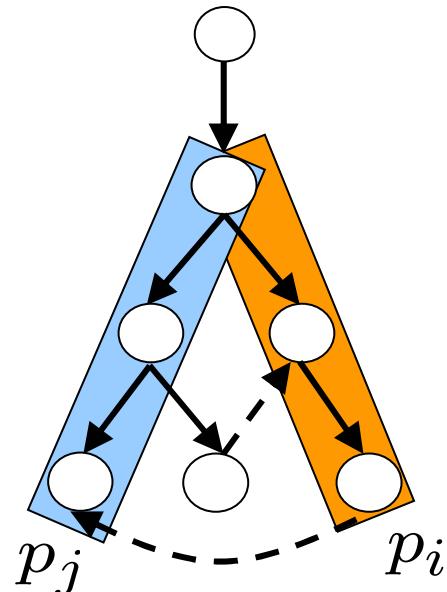


Tree Parameterization

- Construct a spanning tree from the graph
- The mapping function between the poses and the parameters is:

$$x_i = p_i \ominus p_{\text{parent}(i)} \quad X_i = P_{\text{parent}(i)}^{-1} P_i$$

- Error of a constraint in the new parameterization.

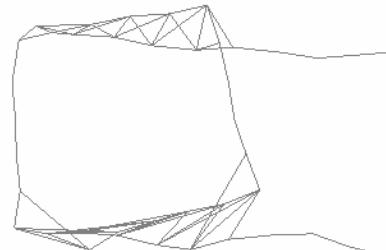
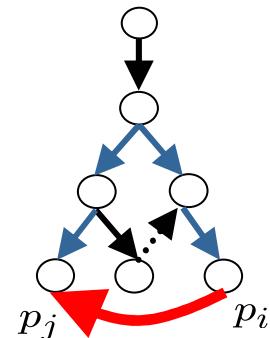
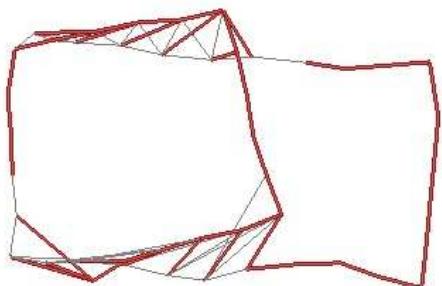


$$E_{ij} = \Delta_{ij}^{-1} \cdot \text{UpChain}^{-1} \cdot \text{DownChain}$$

Only variables in the path of a constraint are involved in the update.

Stochastic Gradient Descent using the Tree Parameterization

- Using a tree parameterization, we decompose the problem in many small sub-problems which are either:
 - constraints on the tree ("open loop")
 - constraints not in the tree ("a loop closure")
- Each SGD equation independently solves one sub-problem at a time
- The solutions are integrated via the learning rate



Computation of the Update Step

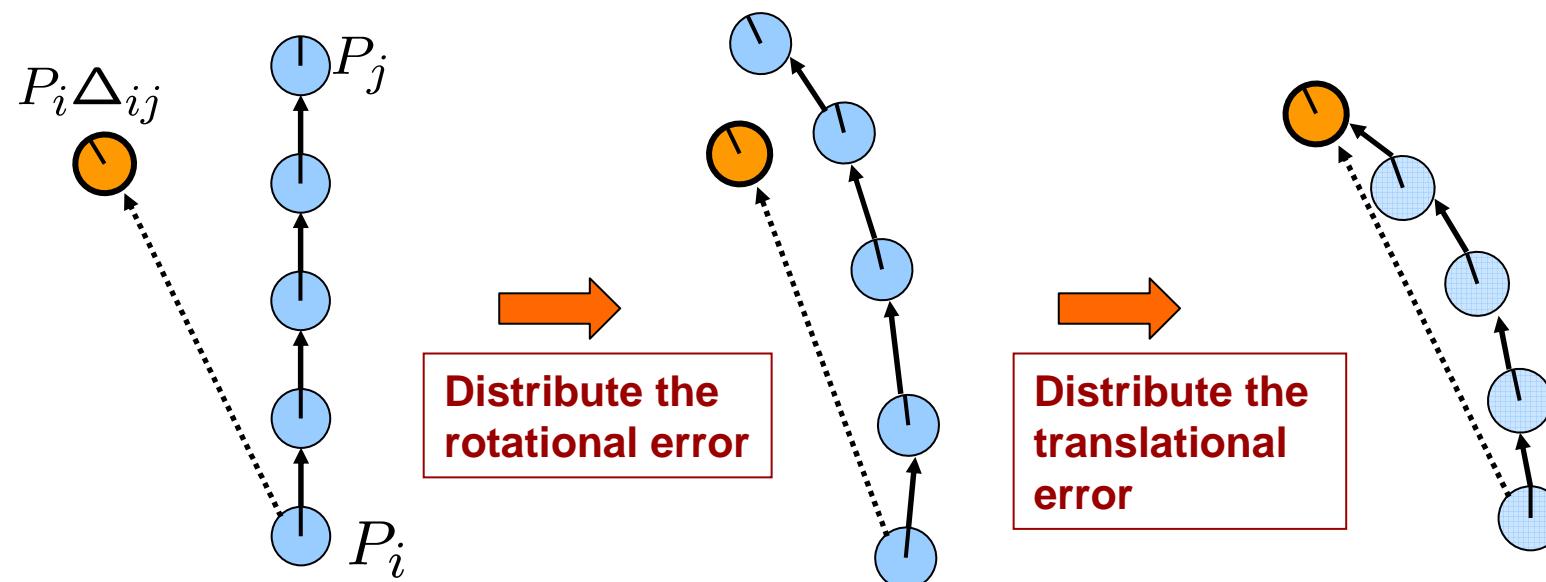
- 3D rotations lead to a highly nonlinear system.
 - Update the poses directly according to the SGD equation may lead to poor convergence.
 - This effect increases with the connectivity of the graph.
- Key idea in the SGD update:

$$\Delta \mathbf{x} = \lambda \cdot \mathbf{H}^{-1} J_{ij}^T \Omega_{ij} r_{ij}$$

Distribute a fraction of the residual along the parameters so that the error of that constraint is reduced.

Computation of the Update Step

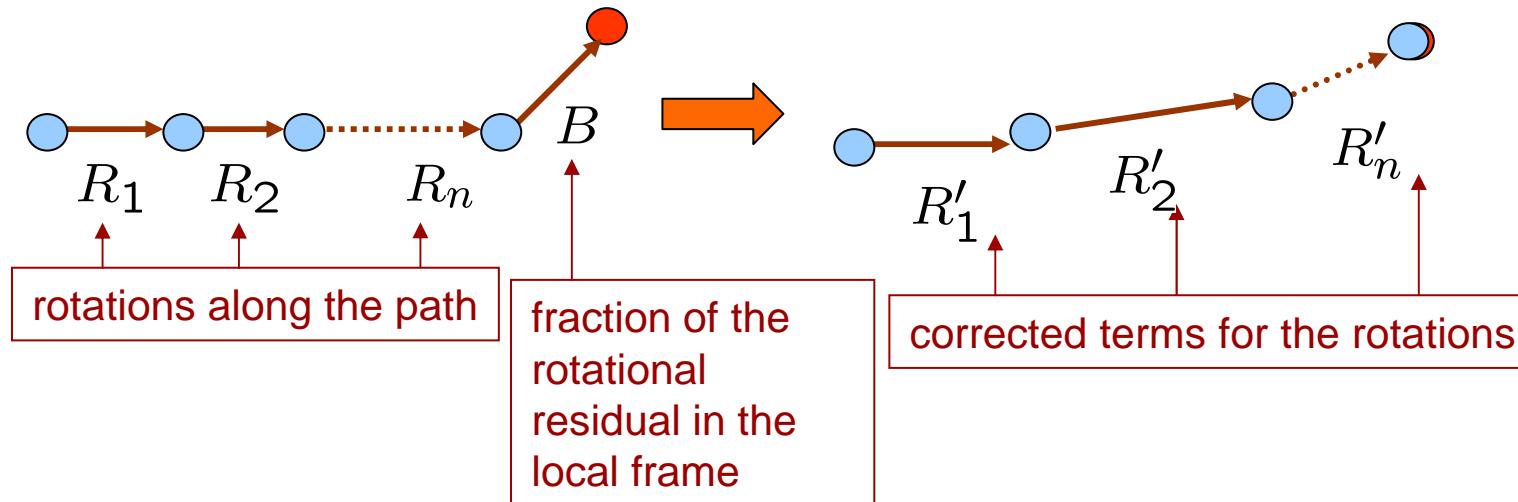
Update in the “spirit” of the SGD:
Smoothly deform the path along the
constraints so that the error is reduced.



Distribution of the Rotational Error

- In 3D the rotational error cannot be simply added to the parameters because the rotations are not commutative.
- Our goal is to find a set of **incremental** rotations so that the following equality holds:

$$R_1 R_2 \cdots R_n B = R'_1 R'_2 \cdots R'_n$$



Distributing the Rotational Residual

- Assume that the first node is the reference frame
- We want a correcting rotation with **a single axis**
- Let A_i be the orientation of the i-th node in the global reference frame

$$A'_n = A_n B = Q A_n$$

with a decomposition of the rotational residual into a chain of incremental rotations obtained by spherical linear interpolation (slerp)

$$Q = Q_1 Q_2 \cdots Q_n$$

$$Q_k = \text{slerp}(Q, u_{k-1})^T \text{slerp}(Q, u_k) \quad u \in [0 \dots \lambda]$$

- Slerp has been designed for 3d animation: constant speed motion along a circle arc

What is the SLERP?

- SLERP = Spherical LinEar inteRPolation
- Introduced by Ken Shoemake for interpolations in 3D animations
- Constant speed motion along a circle arc with unit radius
- Properties:

$$\mathcal{R}' := \text{slerp}(\mathcal{R}, u)$$

$$\text{axisOf}(\mathcal{R}') = \text{axisOf}(\mathcal{R})$$

$$\text{angleOf}(\mathcal{R}') = u \cdot \text{angleOf}(\mathcal{R})$$

Distributing the Rotational Residual

- To obtain the rotations, we recursively solve:

$$R'_1 = Q_1 R_1$$

$$R'_2 = (Q_1 R_1)^T Q_{1:2} R_{1:2}$$

⋮

$$R'_k = [(R_{1:k-1})^T Q_k R_{1:k-1}] R_k$$

- With this update rule, it can be shown that the change in each rotational residual is bounded

$$\Delta r'_{k,k-1} \leq |\text{angleOf}(Q_k)|$$

- This bounds a potentially introduced error at node k when correcting a chain of poses including k

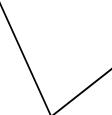
How to Determine u_k ?

- The values of u_k describe the relative distribution of the error along the chain

$$Q_k = \text{slerp}(Q, u_{k-1})^T \text{slerp}(Q, u_k) \quad u \in [0 \dots \lambda]$$

- Here, we need to consider the uncertainty of the constraints

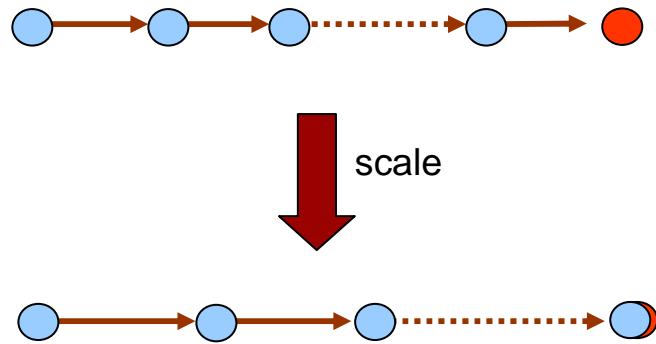
$$u_k = \min \left(1, \lambda |\mathcal{P}_{ij}| \right) \left[\sum_{m \in \mathcal{P}_{ij} \wedge m \leq k} d_m^{-1} \right] \left[\sum_{m \in \mathcal{P}_{ij}} d_m^{-1} \right]^{-1}$$


 $d_m = \sum_{\substack{\text{all constraints connecting } m \\ \langle l, m \rangle}} \min [\text{eigen}(\Omega_{lm})]$

- This assumes roughly spherical covariances!

Distributing the Translation Part of the Residual

- That is trivial
- Just scale the x, y, z dimension (with covariances)



Summary of the Algorithm

- Decompose the problem according to the tree parameterization
- Loop
 - Select a constraint
 - Randomly
 - Alternative: sample inverse proportional to the number of nodes involved in the update
 - Compute the nodes involved in the update
 - Nodes according to the parameterization tree
 - Reduce the error for this sub-problem
 - Reduce the rotational error (slerp)
 - Reduce the translational error

Complexity

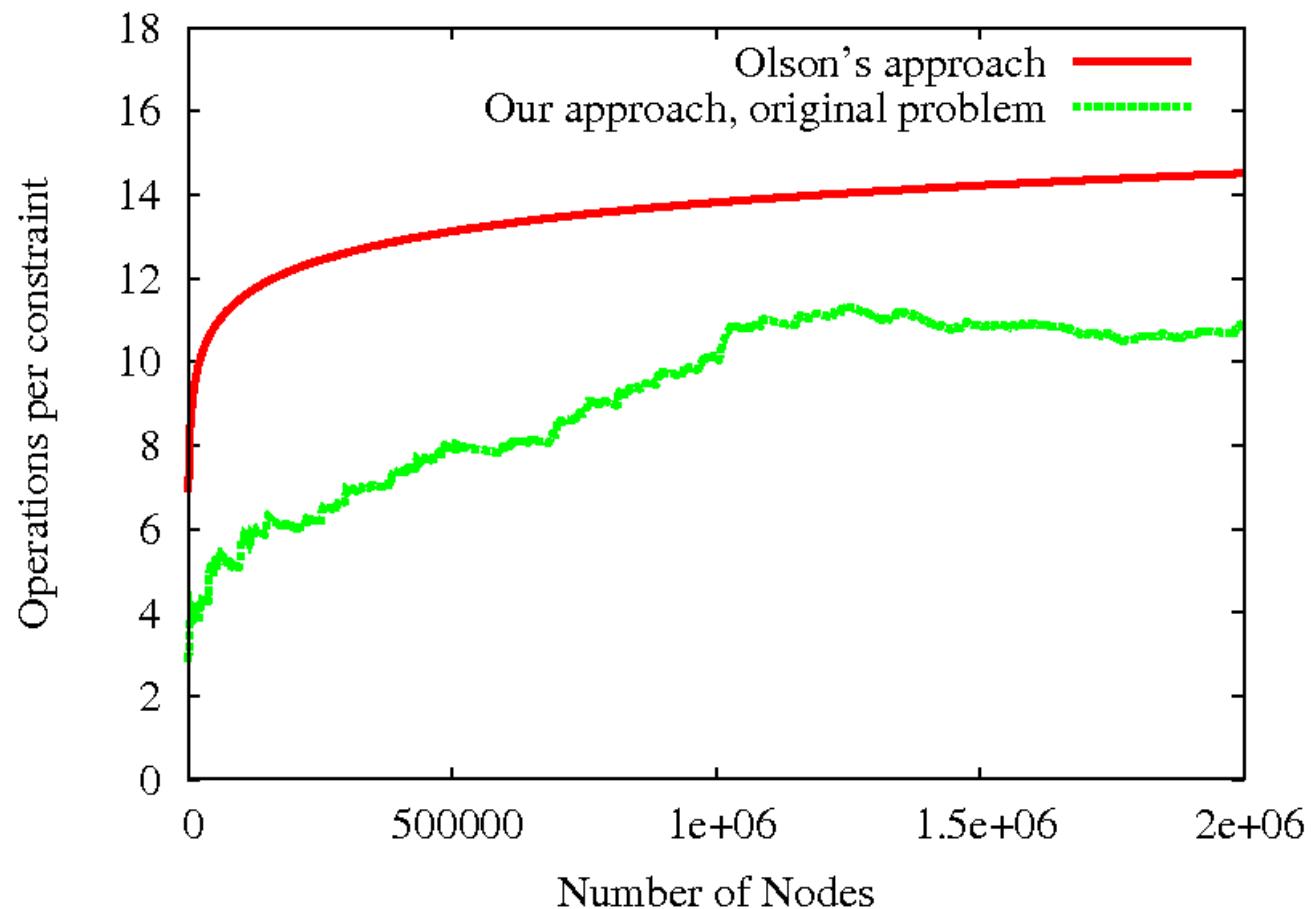
- In each iteration, the approach considers all constraints
- Each constraint optimization step requires to update a set of nodes (on average: the average “path length according to the tree)
- This results in a complexity per iteration of

$$\mathcal{O}(M \cdot l)$$

#constraints

avg. path length
(parameterization tree)

Cost of a Constraint Update



$$\approx \mathcal{O}(M \cdot \log(N))$$

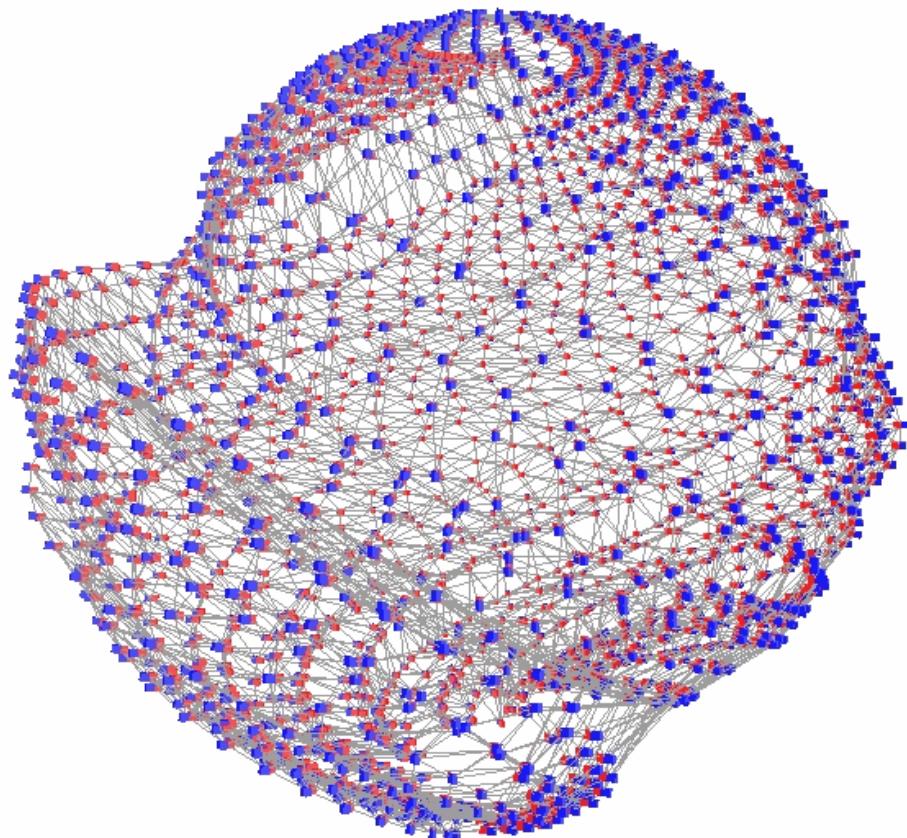
Node Reduction

- The complexity grows with the length of the trajectory
- That is bad for life-long learning!
- Idea: Combine constraints between nodes if the robot is well-localized

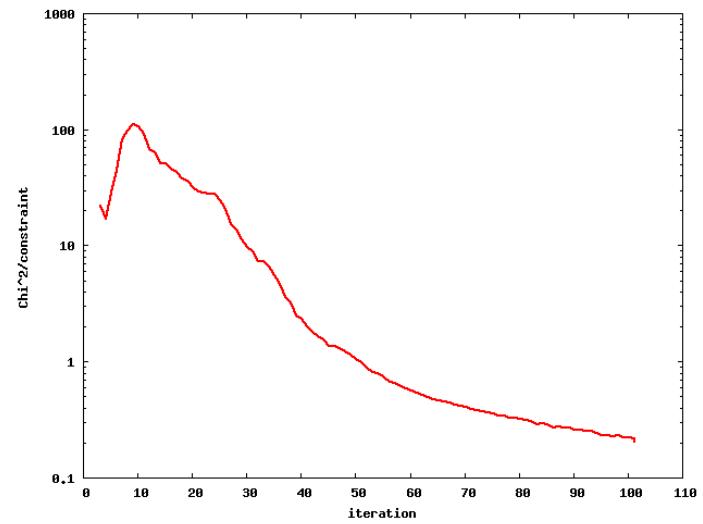
$$\begin{aligned}\Omega_{ij} &= \Omega_{ij}^{(1)} + \Omega_{ij}^{(2)} \\ \delta_{ij} &= \Omega_{ij}^{-1}(\Omega_{ij}^{(1)}\delta_{ij}^{(1)} + \Omega_{ij}^{(2)}\delta_{ij}^{(2)})\end{aligned}$$

- This is similar to adding a rigid constraints
- Valid approximation if the robot is well localized
- Complexity depends only on the size if the environment, not the length of the trajectory

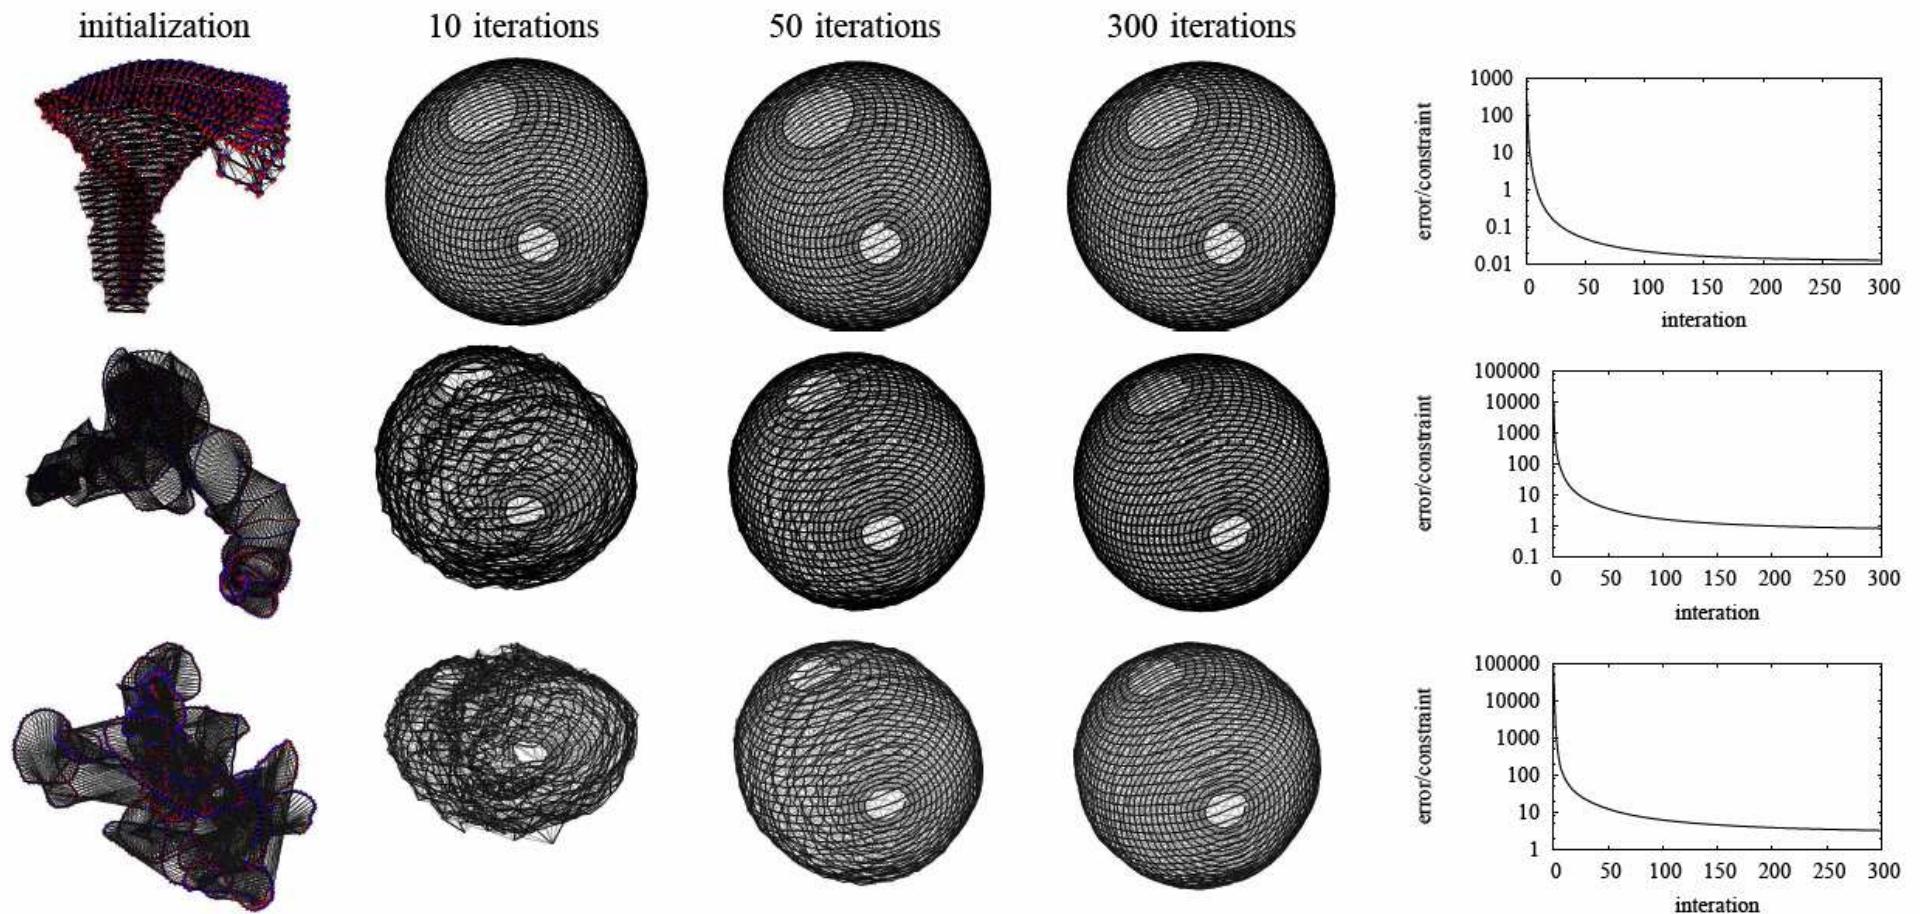
Simulated Experiment



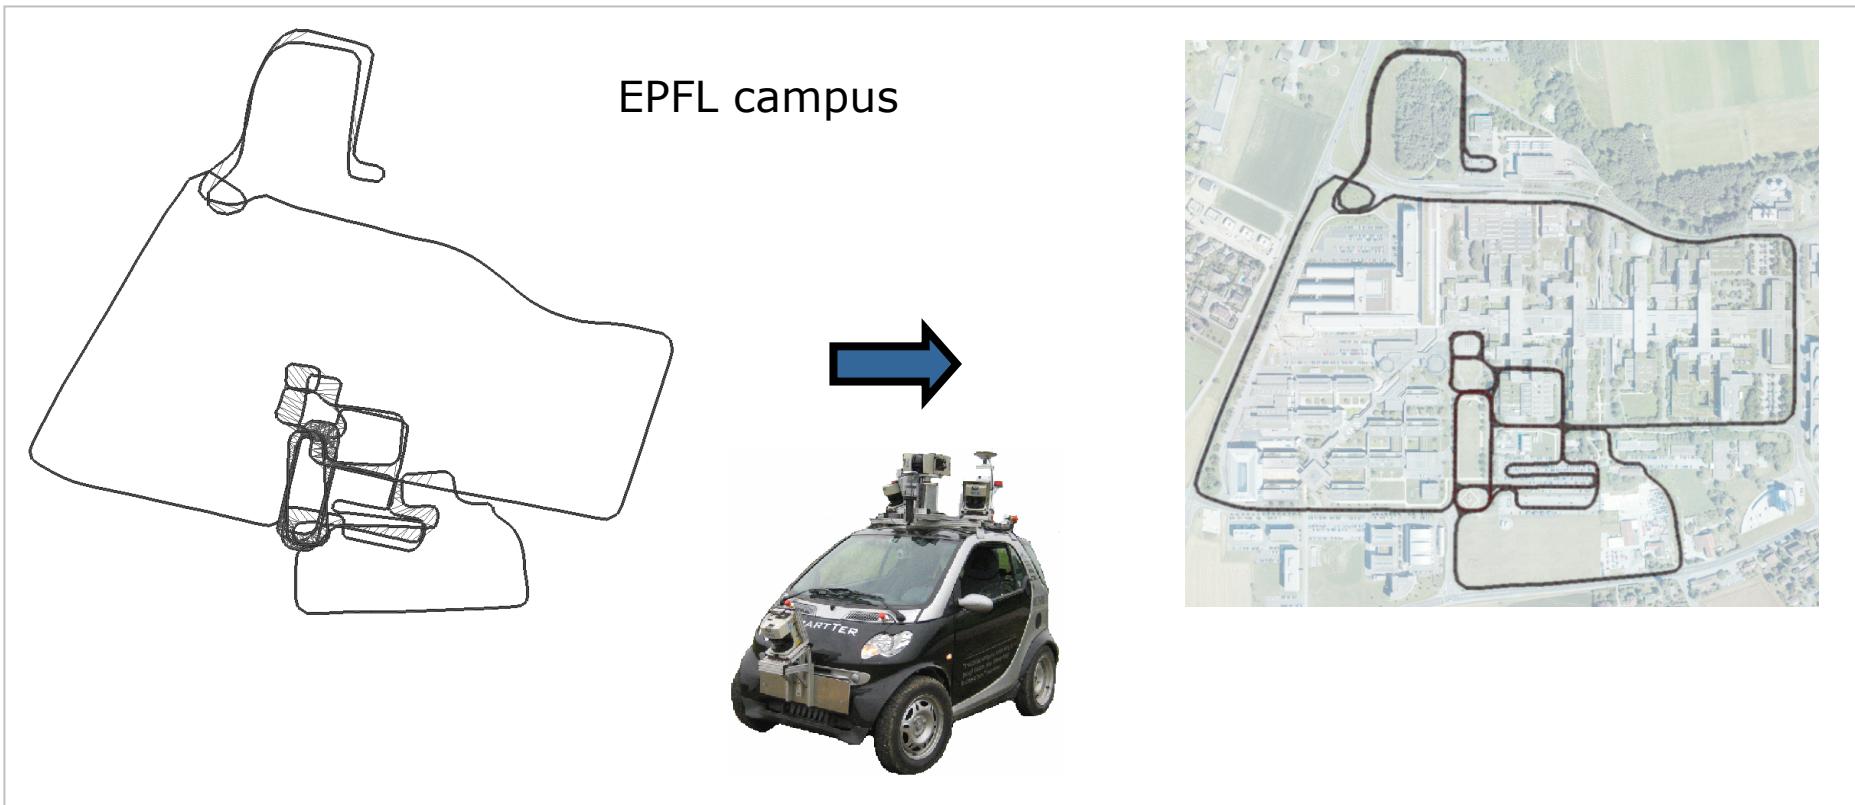
- Highly connected graph
- Poor initial guess
- LU & friends fail
- 2200 nodes
- 8600 constraints



Spheres with Different Noise

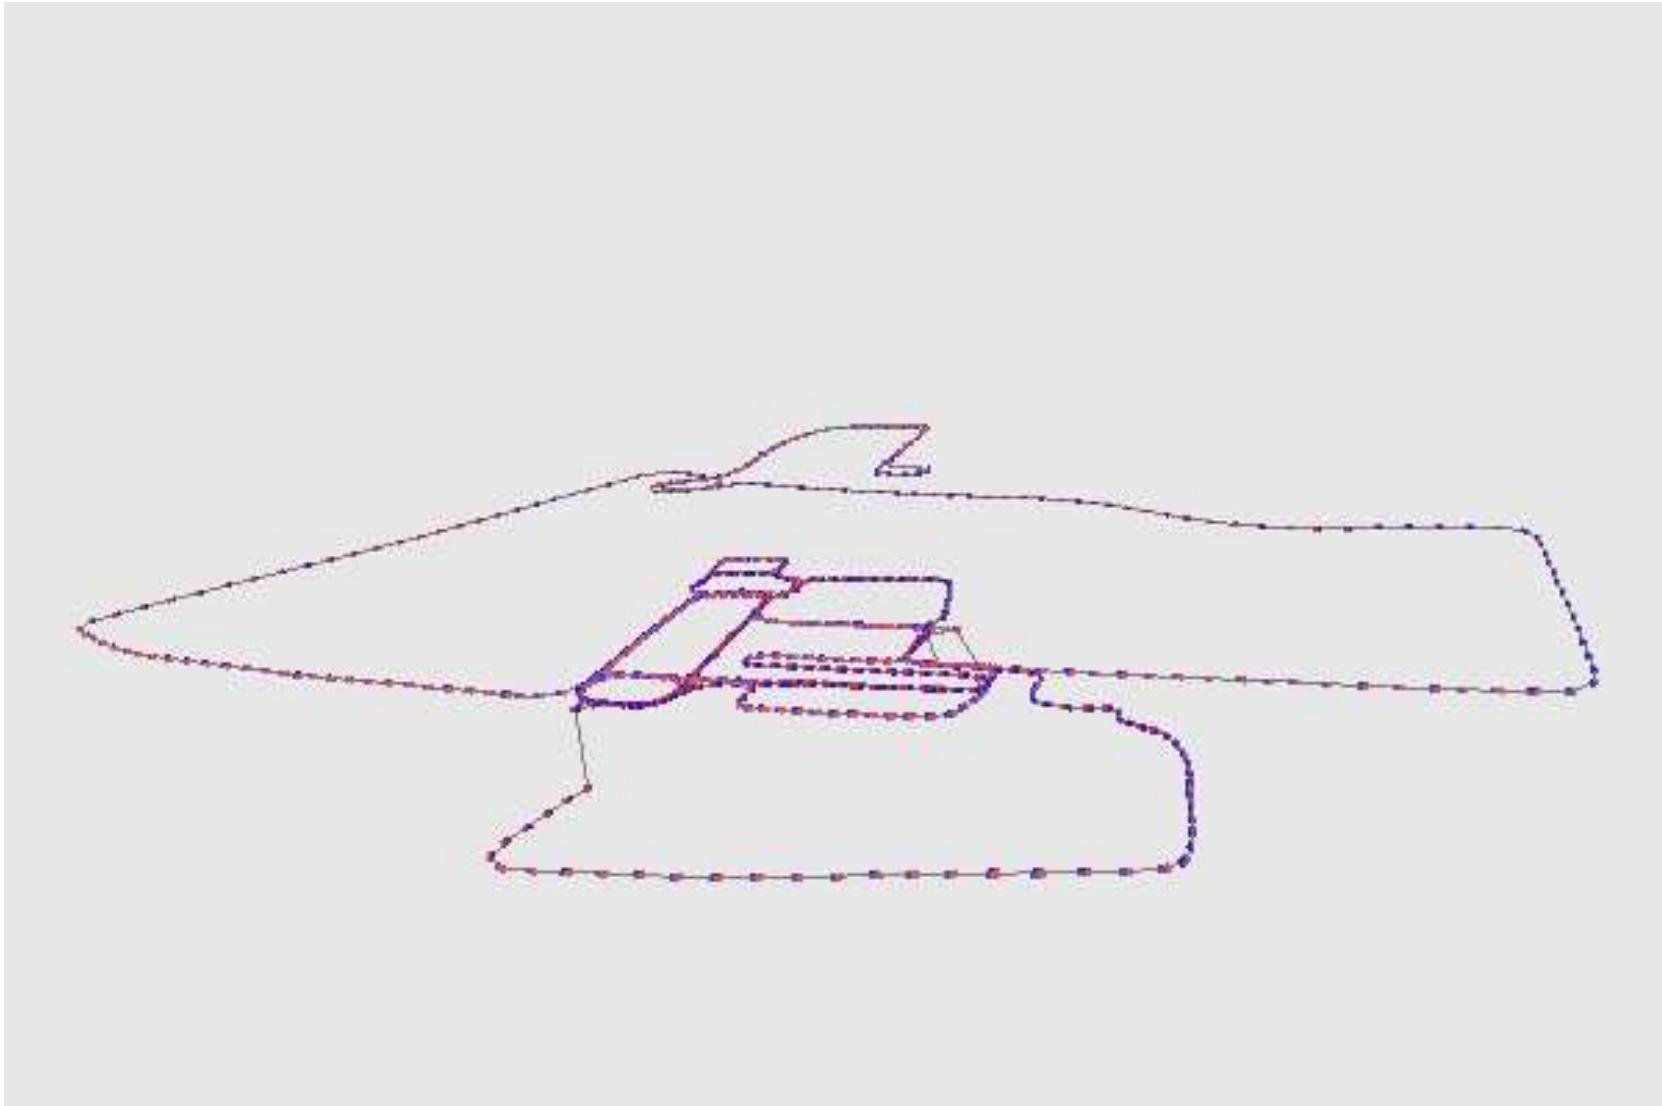


Mapping the EPFL Campus



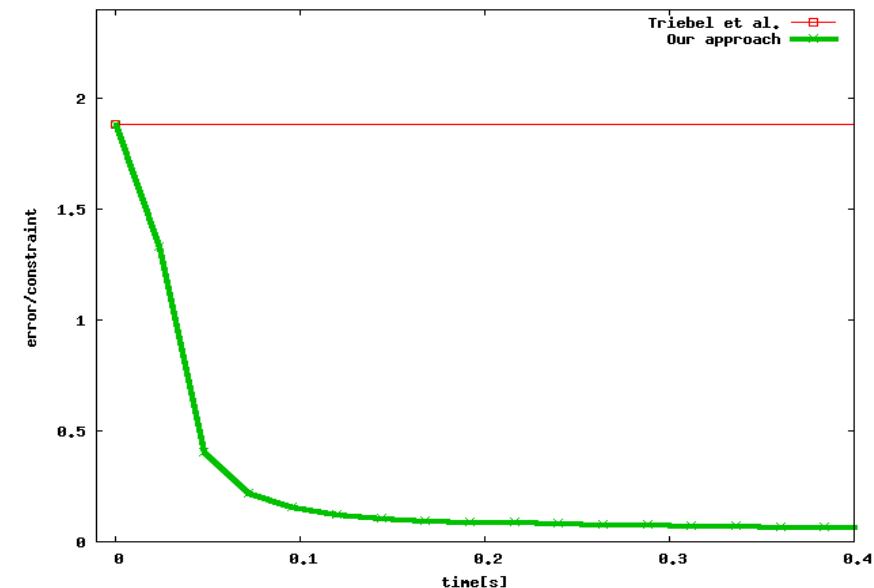
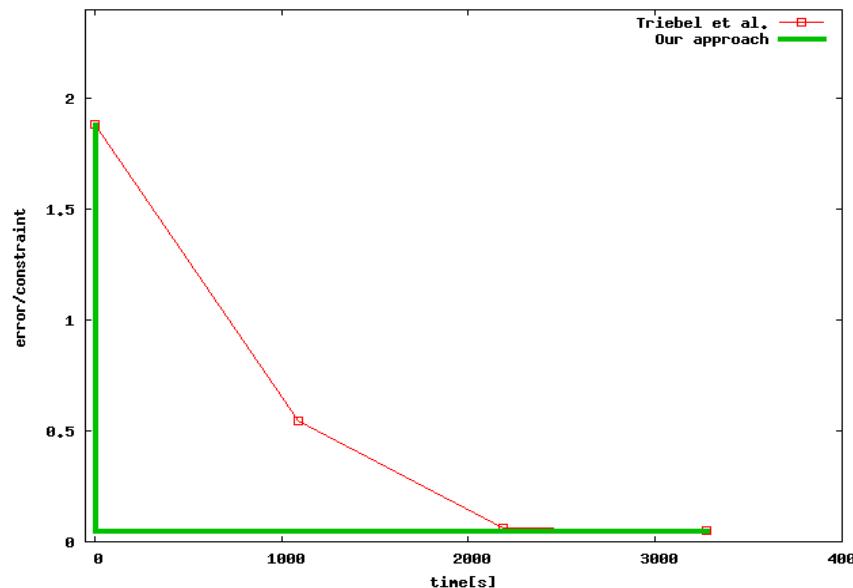
- 10km long trajectory and 3D lasers recorded with a car
- Problem not easily tractable by most standard optimizers

Mapping the EPFL Campus



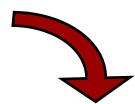
Comparison with Standard Approaches (LU Decomposition)

- Tractable subset of the EPFL dataset
- Optimization carried out in less than one second.
- The approach is so fast that in typical applications one can run it while incrementally constructing the graph.



TORO vs. Olson's Approach

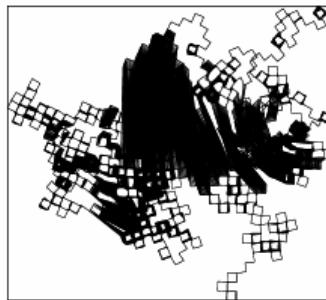
Olson's approach



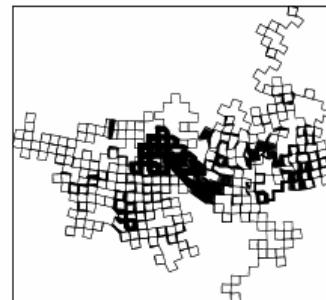
1 iteration



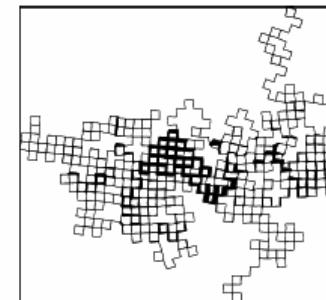
10 iterations



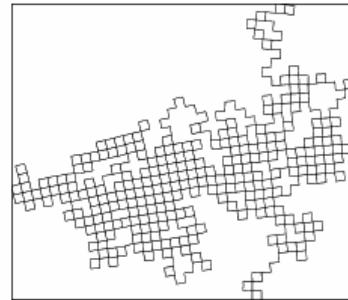
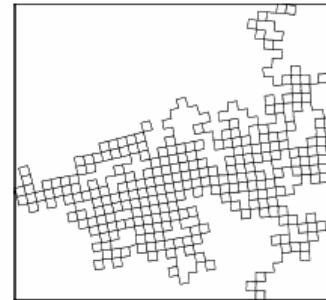
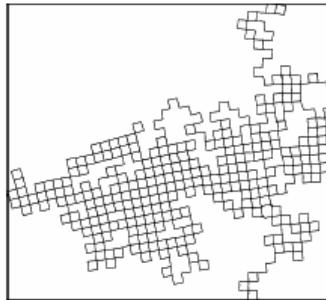
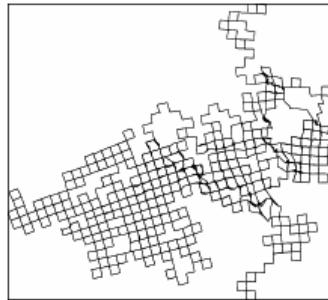
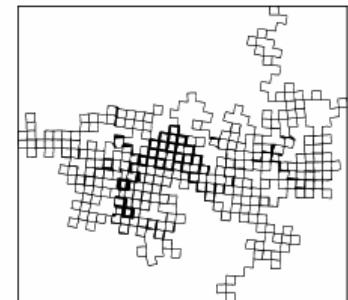
50 iterations



100 iterations

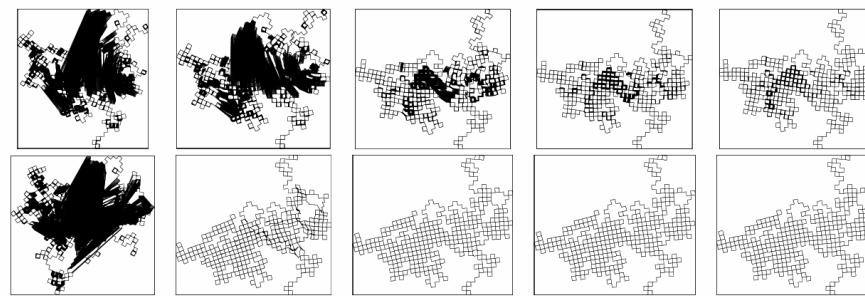
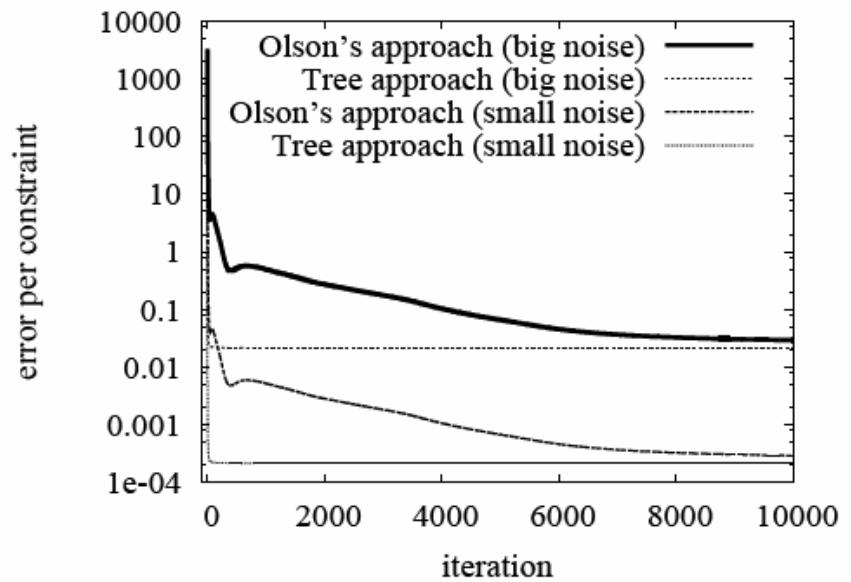
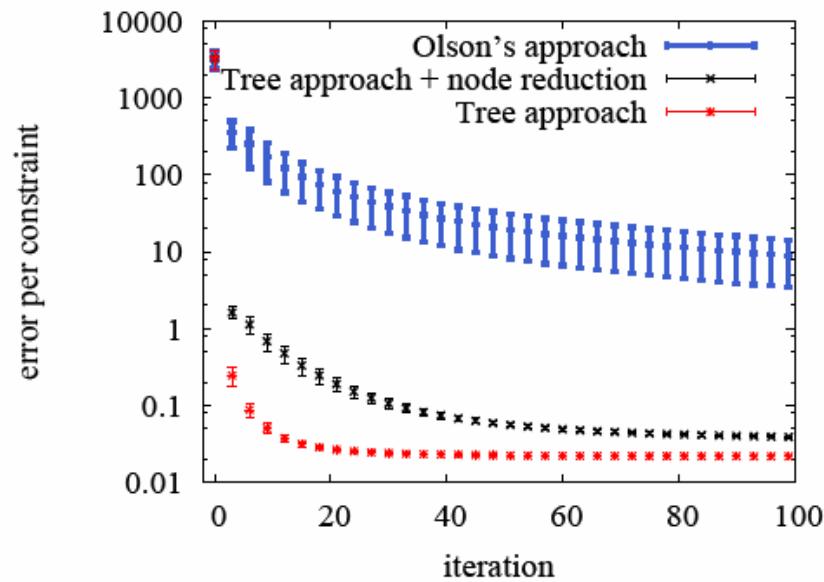


300 iterations

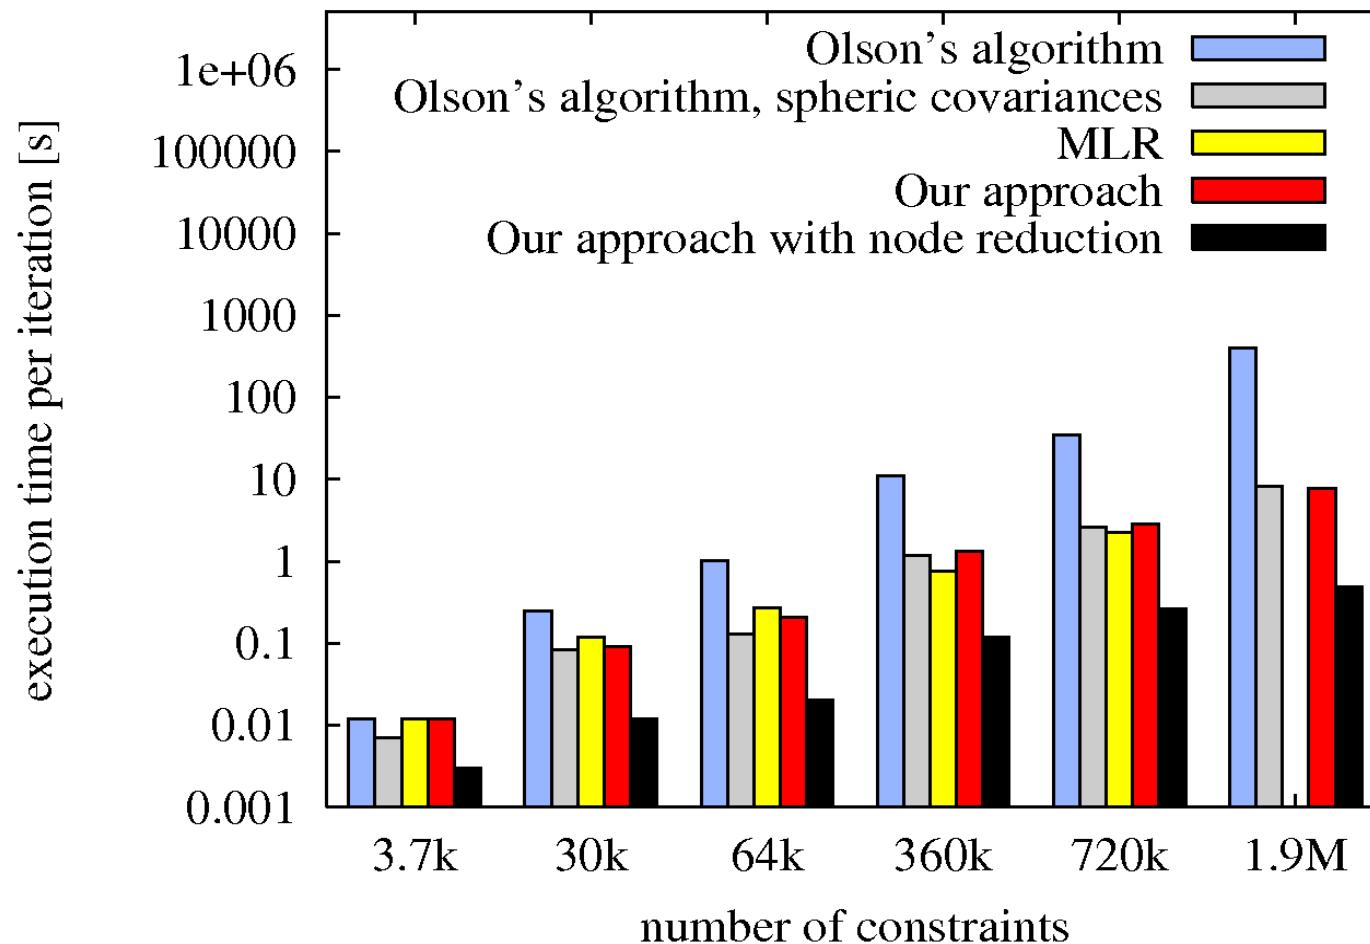


TORO

TORO vs. Olson's Approach

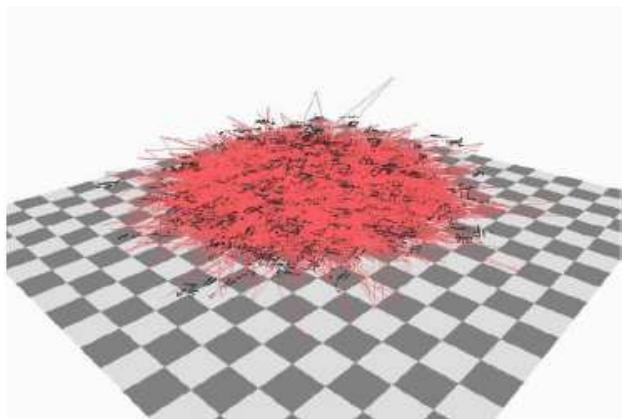


Time Comparison (2D)

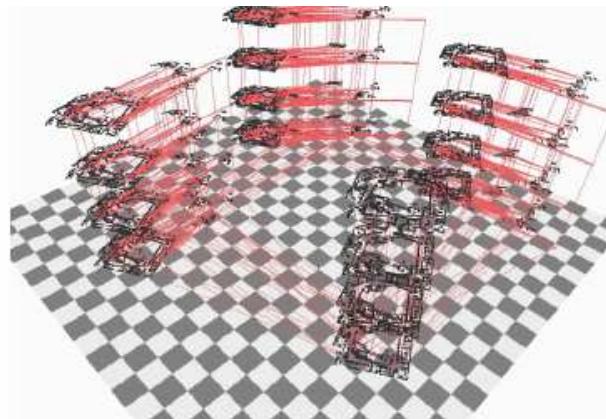


Robust to the Initial Guess

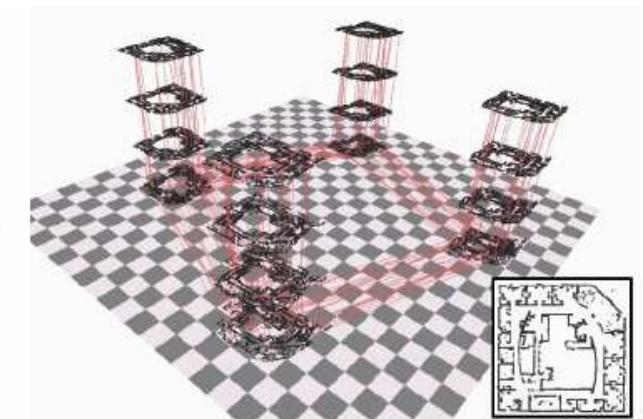
- Random initial guess
- Intel datatset as the basis for 16 floors distributed over 4 towers



initial configuration



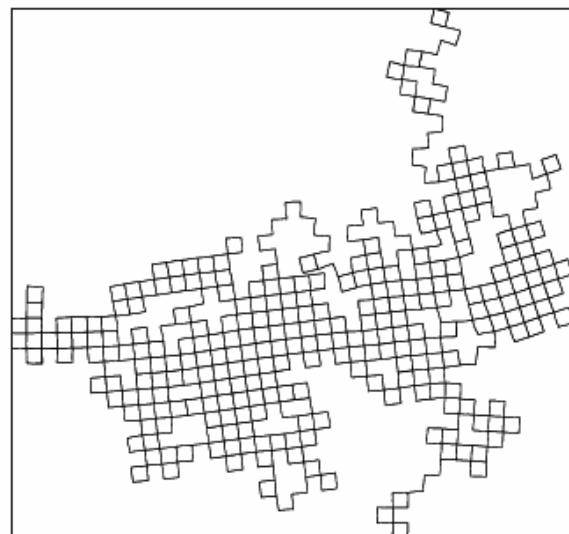
intermediate result



final result (50 iterations)

Informal Comparison to MLR

- More robust under the initial guess
- Much faster than MLR
- If MLR converges, it gives slightly lower error values



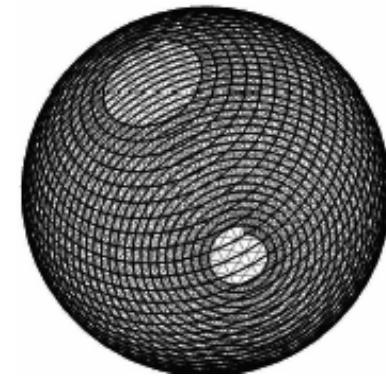
with a good initial guess



with a bad initial guess

Informal Comparison to (i)SAM

- More stable than SAM
- Results similar to iSAM
- Faster than SAM/iSAM



noise level	SAM (batch)	SAM (incremental)	Our method (batch)
$\sigma = 0.05$	119 s	not tested (see batch)	20 s (100 iterations)
$\sigma = 0.1$	diverged	270 s (optimized each 100 nodes)	40 s (200 iterations)
$\sigma = 0.2$	diverged	510 s (optimized each 50 nodes)	50 s (250 iterations)

TORO Summary

- Extremely robust to bad initial network configurations
- One of the most efficient techniques for ML map estimates that is available (~01/2008)
- Works in 2D and 3D
- Scales up to millions of constraints
- Allows for multi-vehicle mapping (with minor modifications)
- Used by different research groups/companies worldwide
(U. Oxford, ETH Zurich, La Sapienza Rome, Willow Garage, KUKA, ...)
- Download from <http://www.openslam.org/toro.html>



Drawbacks of TORO

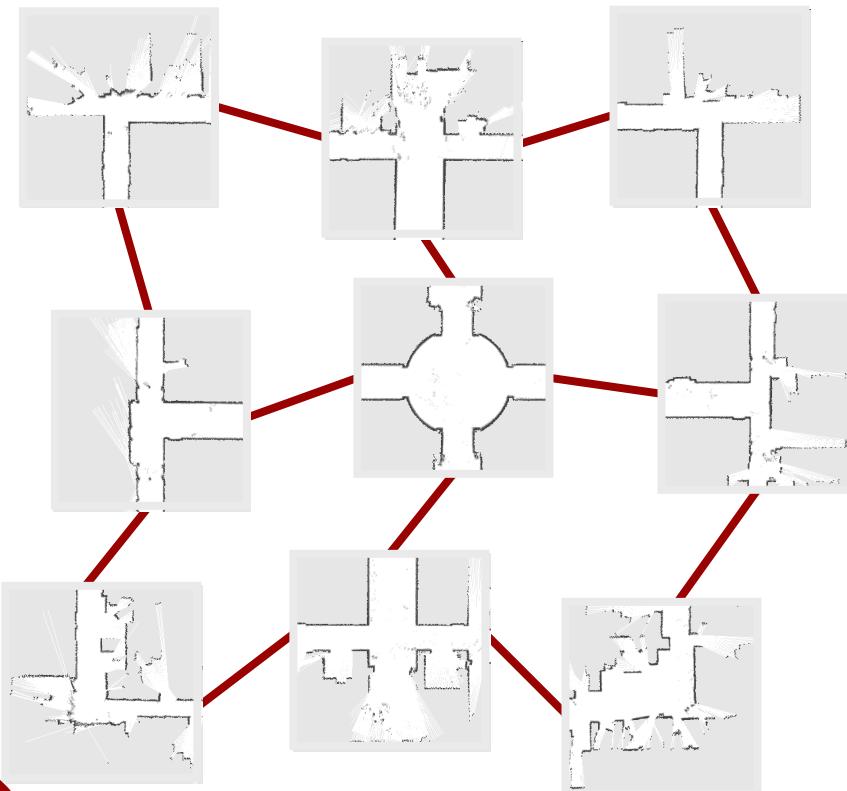
- The slerp-based update rule optimizes rotations and translations separately.
- It assume **roughly spherical covariance ellipses**.
- It is a maximum likelihood technique.
No covariance estimates!
- Approach of Tipaldi et al. accurately estimates the covariances after convergence [Tipaldi et al., 2007]

Data Association

- So far we explained how to compute the mean of the distribution given the data associations.
- However, to determine the data associations, we need to know the covariance matrices of the nodes.
- Standard approaches include:
 - Matrix inversion
 - Loopy belief propagation
 - Belief propagation on a spanning tree
 - **Loopy intersection propagation**

Graphical SLAM as a GMRF

- Factor the distribution
 - local** potentials
 - pairwise** potentials



$$p(x) = \frac{1}{Z} \prod_{i=1}^n \phi_i(x_i) \prod_{j=i+1}^n \underline{\phi_{i,j}(x_i, x_j)}$$

Gaussian in canonical form

Belief Propagation

- Inference by local message passing
- Iterative process
 - **Collect** messages

$$\underline{m_i^{(t)} = \eta_i + \sum_{j \in \mathcal{N}_i} m_{ji}^{(t-1)}}$$

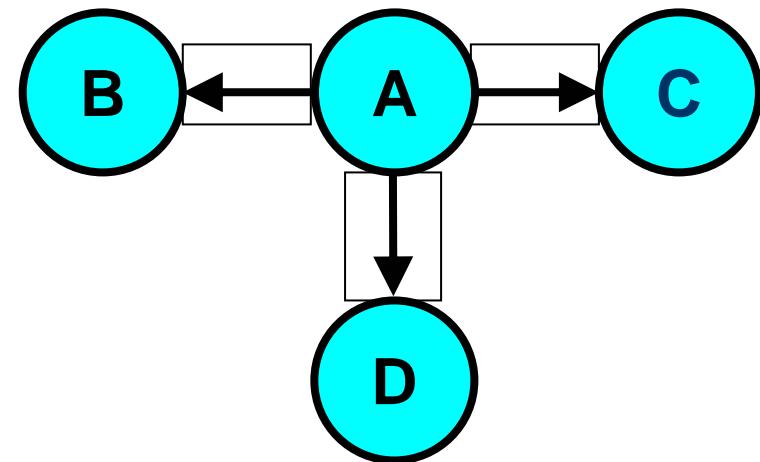
$$M_i^{(t)} = \Omega_i + \sum_{j \in \mathcal{N}_i} M_{ji}^{(t-1)}$$

- **Send** messages

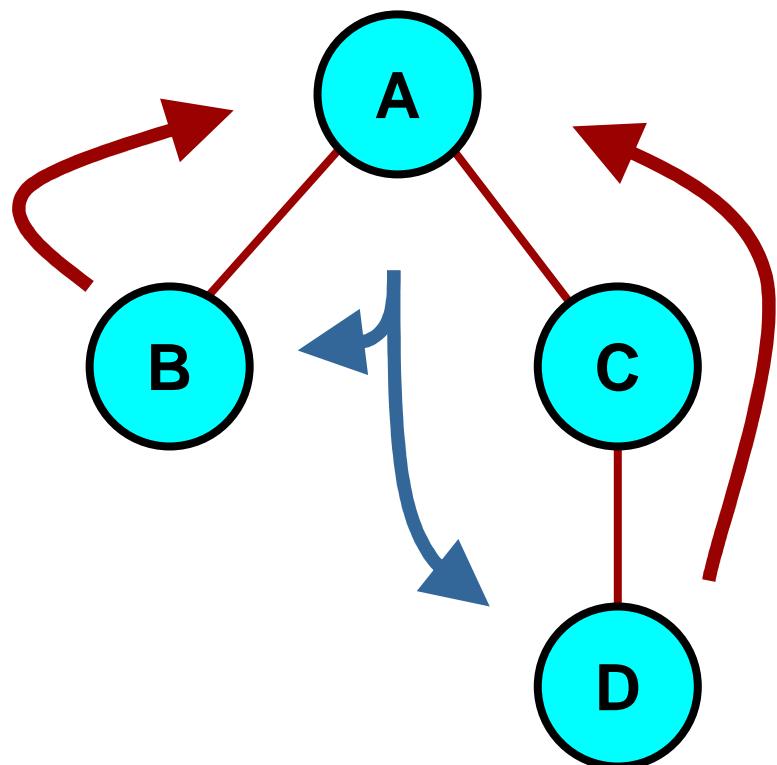
$$\underline{m_{ij}^{(t)} = \eta_{ij}^j - \Omega_{ij}^{[ji]} \left(\Omega_{ij}^{[ii]} + M_i^{(t)} - M_{ji}^{(t-1)} \right)^{-1} \left(\eta_{ij}^i + m_i^{(t)} - m_{ji}^{(t-1)} \right)}$$

$$M_{ij}^{(t)} = \boxed{\Omega_{ij}^{[jj]} - \Omega_{ij}^{[ji]} \left(\Omega_{ij}^{[ii]} + M_i^{(t)} - M_{ji}^{(t-1)} \right)^{-1} \Omega_{ij}^{[ij]}}$$

Ignore the math!

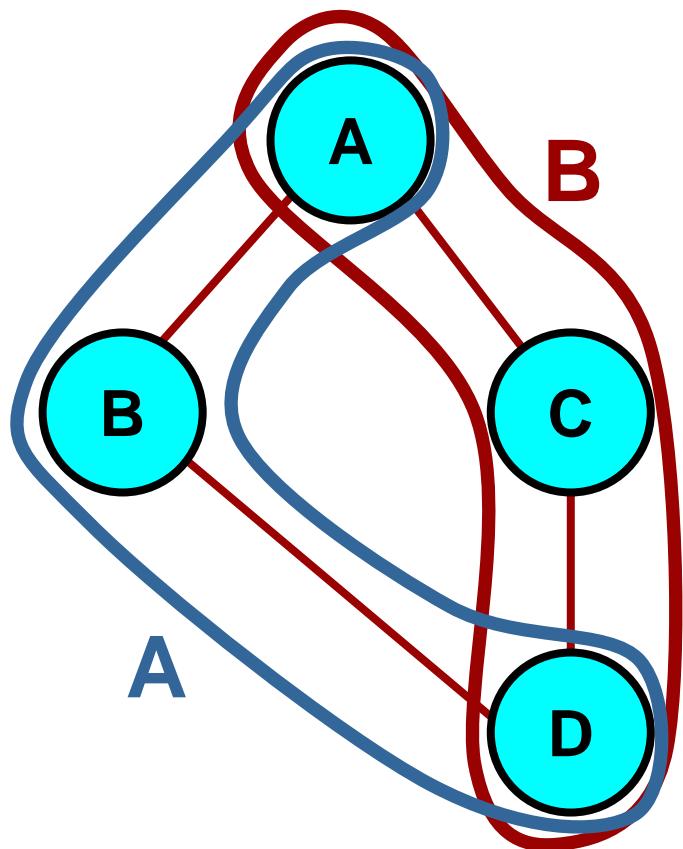


Belief Propagation - Trees



- Exact inference
- Message passing
- Two iterations
 - From leaves to root:
variable elimination
 - From root to leaves:
back substitution

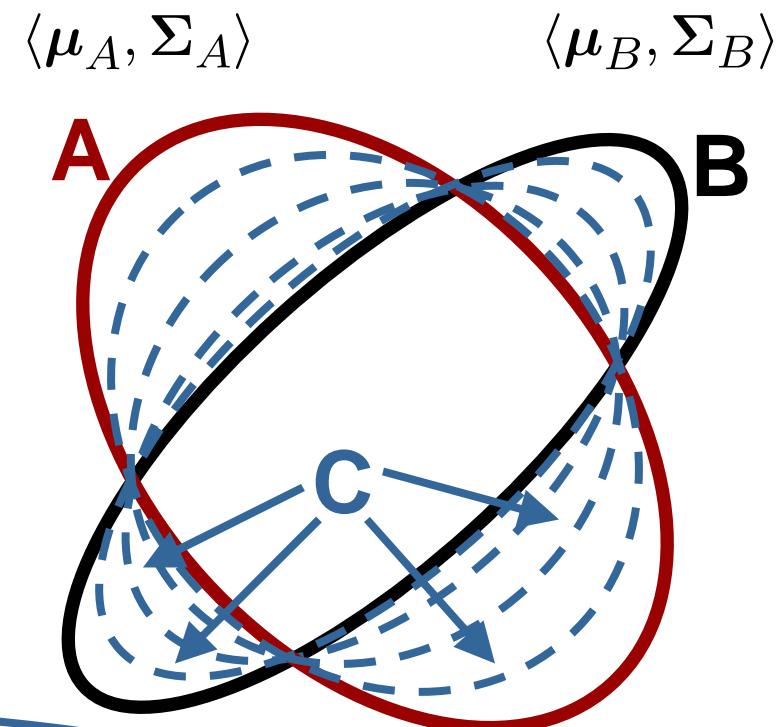
Belief Propagation - loops



- Approximation
- Multiple paths
- Overconfidence
 - Correlations between path **A** and path **B**
- How to integrate information at **D**?

Covariance Intersection

- Fusion rule for unknown correlations
- Combine **A** and **B** to obtain **C**



$$\Sigma_C = (\omega \Sigma_A^{-1} + (1 - \omega) \Sigma_B^{-1})^{-1}$$

$$\mu_C = \Sigma_C(\omega \Sigma_A^{-1} \mu_A + (1 - \omega) \Sigma_B^{-1} \mu_B)$$

Loopy Intersection Propagation

■ Key ideas

- Exact inference on a spanning tree of the graph
- Augment the tree with information coming from loops

■ How

- Approximation by means of **cutting matrices**
- Loop information within **local potentials** (priors)

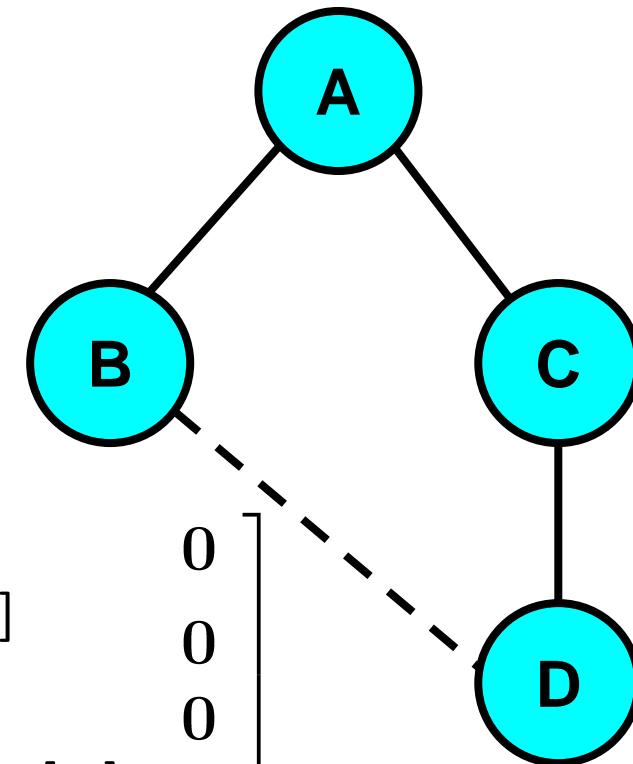
Approximation via Cutting Matrix

- Removal as matrix subtraction

$$\hat{\Omega} = \Omega - K$$

- Regular cutting matrix
- Cut our all off-tree edges

$$K_{BD} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & \Omega_{BD}^{[BB]} - P_{BD}^{[B]} & 0 & \Omega_{BD}^{[BD]} & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & \Omega_{BD}^{[DB]} & 0 & \Omega_{BD}^{[DD]} - P_{BD}^{[D]} & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



Fusing Loops with Spanning Trees

- Estimate **A** and **B**

$$\begin{aligned} E_{BD}^{[D]} &= \Omega_{BD}^{[BB]} - \Omega_{BD}^{[BD]}(M_D + \Omega_{BD}^{[DD]})^{-1}\Omega_{BD}^{[DB]} \\ E_{BD}^{[B]} &= \Omega_{BD}^{[DD]} - \Omega_{BD}^{[DB]}(M_B + \Omega_{BD}^{[BB]})^{-1}\Omega_{BD}^{[BD]} \end{aligned}$$

Ignore the math!

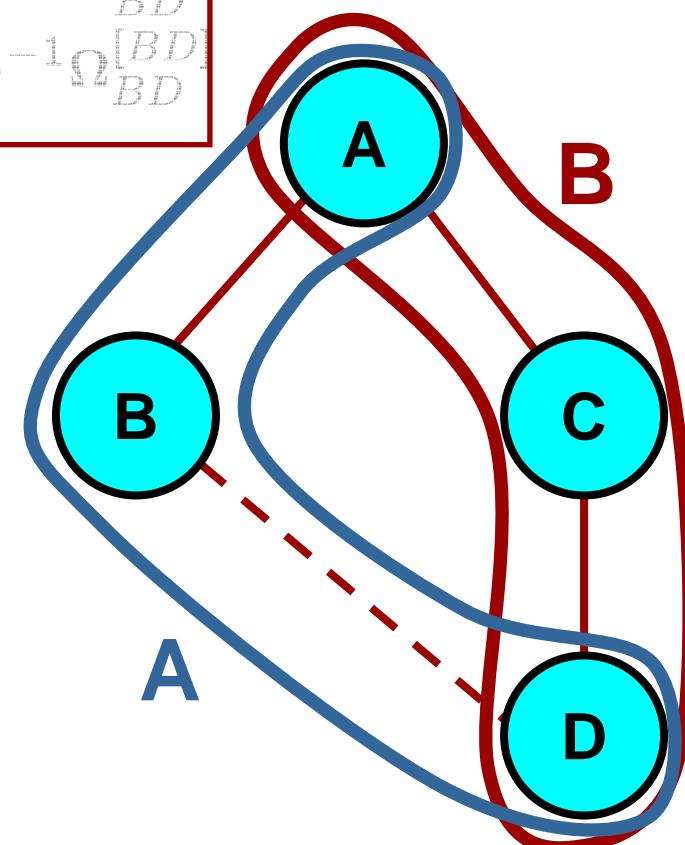
- Fuse the estimates

$$\begin{aligned} \hat{M}_B &= \omega_B M_B + (1 - \omega_B)E_{BD}^{[B]} \\ \hat{M}_D &= \omega_D M_D + (1 - \omega_D)E_{BD}^{[D]} \end{aligned}$$

Ignore the math!

- Compute the priors

$$P_{ij}^{[k]} = \hat{M}_k - M_k$$



LIP – Algorithm

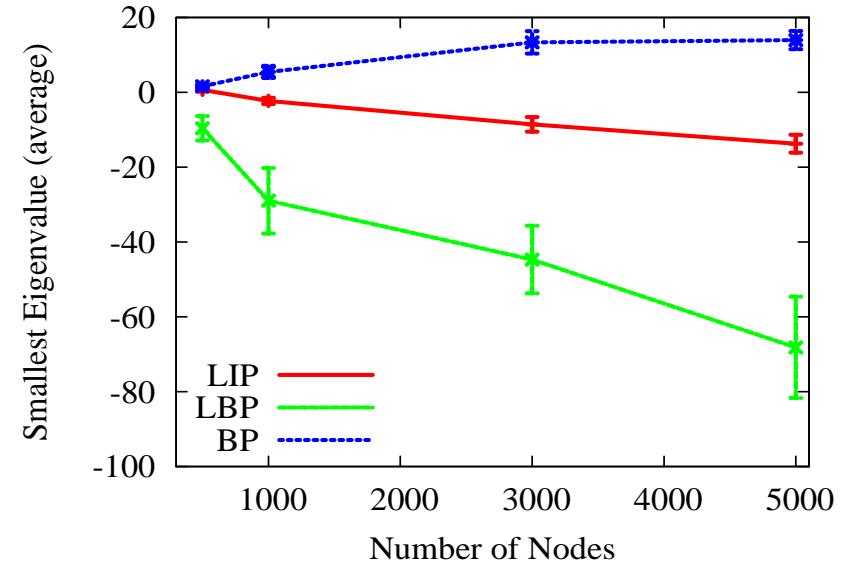
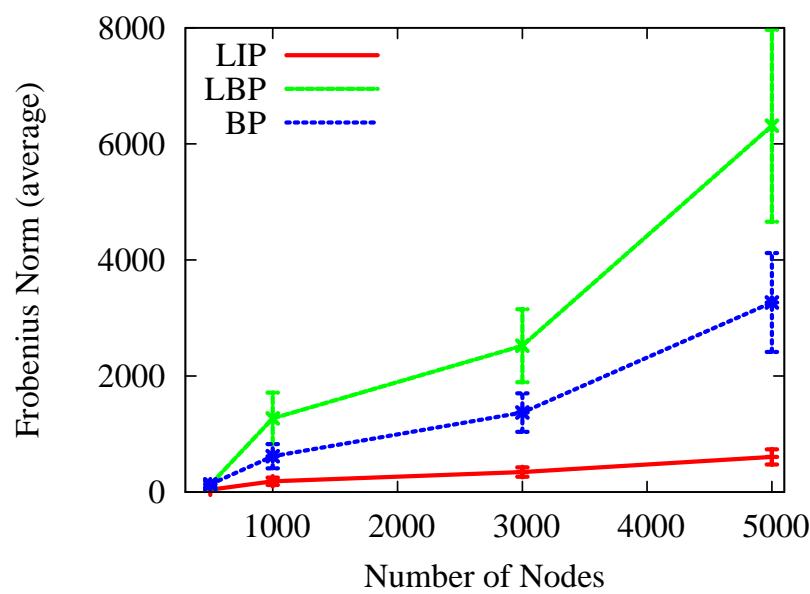
1. Compute a spanning tree
2. Run belief propagation on the tree
3. For every off-tree edge
 1. compute the off-tree estimates,
 2. compute the new priors, and
 3. delete the edge
4. Re-run belief propagation

Experiments – Setup & Metrics

- Simulated data
 - Randomly generated networks of different sizes
- Real data
 - Graph extracted from Intel and ACES dataset from radish
- Approximation error
 - Frobenius norm
- Conservativeness
 - Smallest eigenvalue of matrix difference

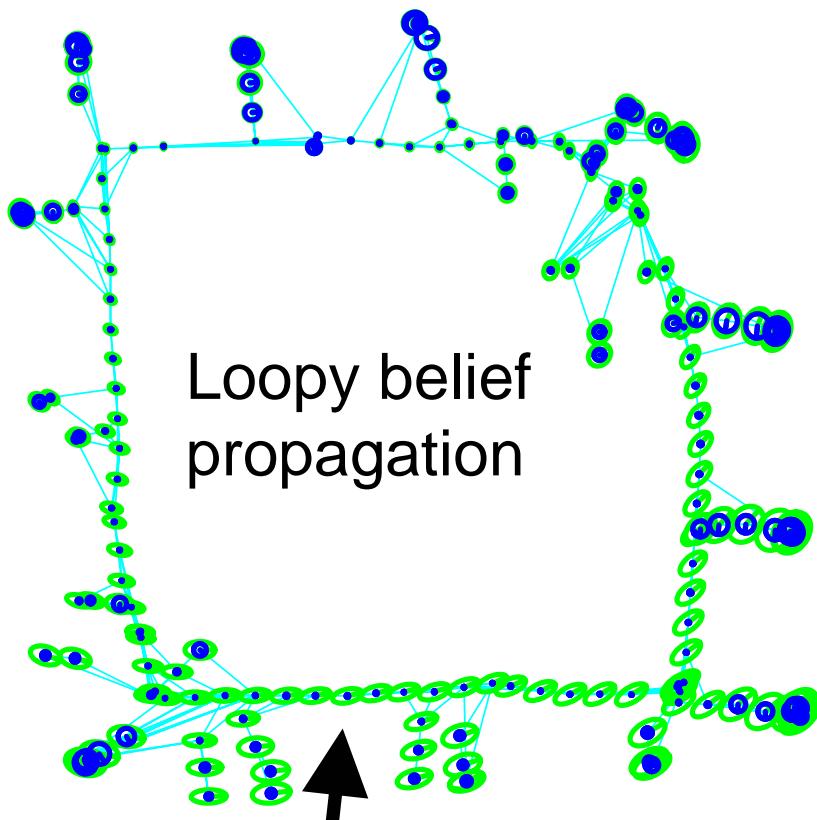
Experiments - Simulated Data

Approximation error

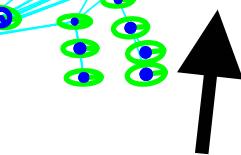


Conservativeness

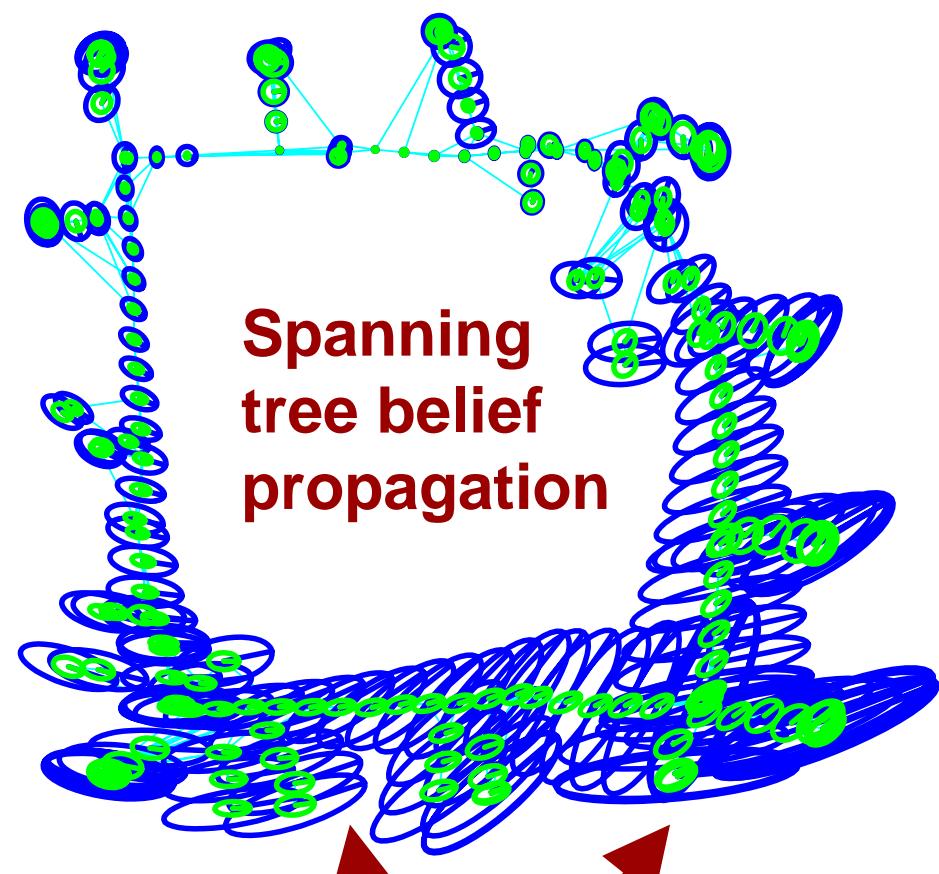
Experiments – Real Data (Intel)



Loopy belief
propagation



Overconfident



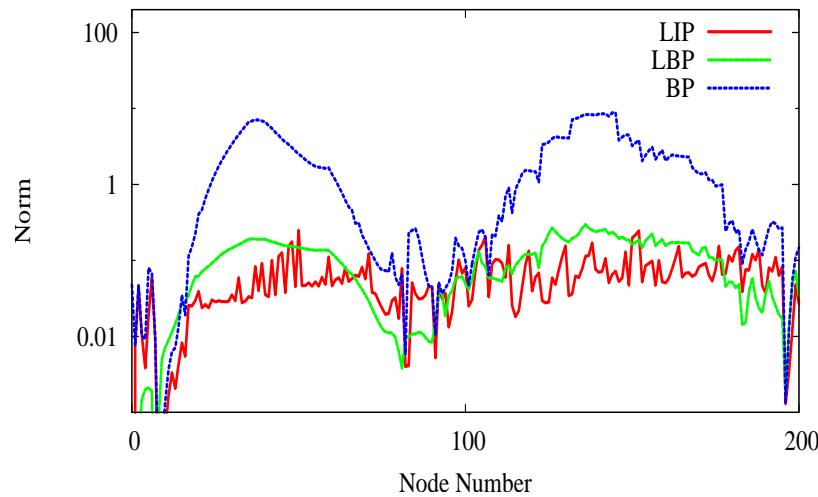
Spanning
tree belief
propagation



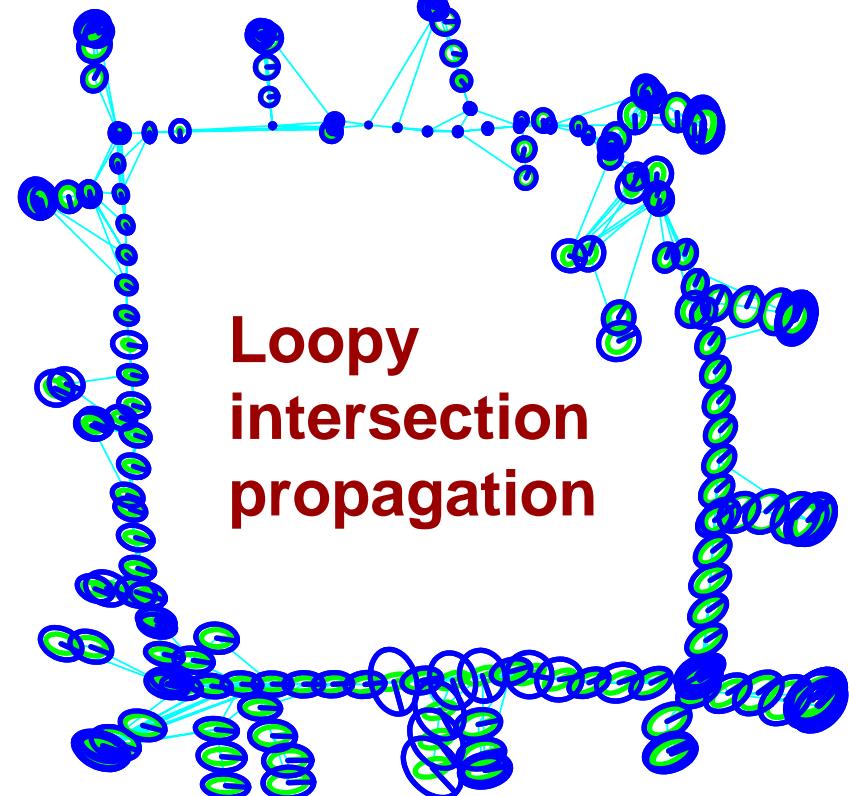
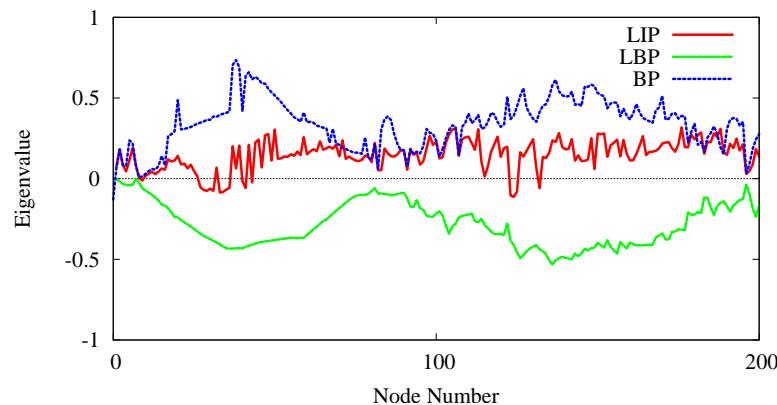
Too conservative

Experiments – Real Data (Intel)

Approximation Error



Conservativeness



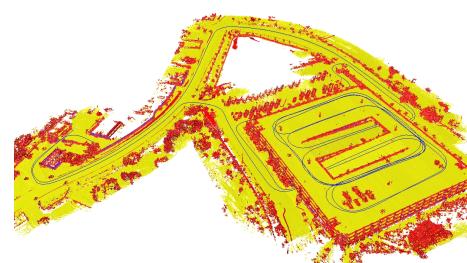
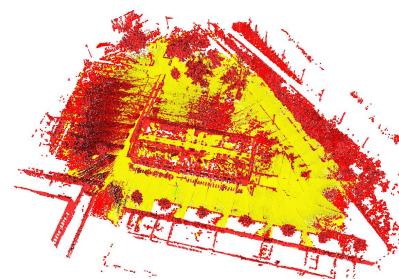
Conclusions – So far...

- TORO - Efficient maximum likelihood algorithm for 2D and 3D graphs of poses
- But no covariance estimates!
- Approach for recovering the covariance matrices via belief propagation and covariance intersection
 - Linear time complexity
 - Tight estimates
 - Generally conservative (not guaranteed!)

SLAM Front-end and Back-end

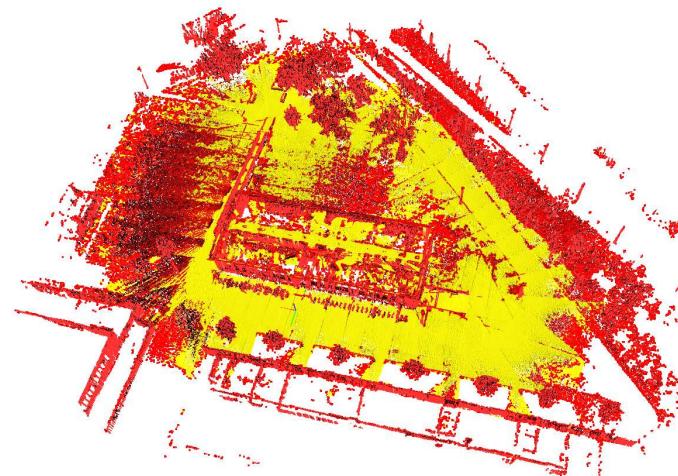
- So far, we talked only about the “SLAM back-end” (how to optimize a network)
- For real world applications, the “front-end” is important as well
- How to interpret the sensor data. This encodes
 - Data association problems
 - Extraction of spatial relations/constraints
- The remainder of this talk will give examples for SLAM front-ends

Three Applications Scenarios that use TORO to Build Maps



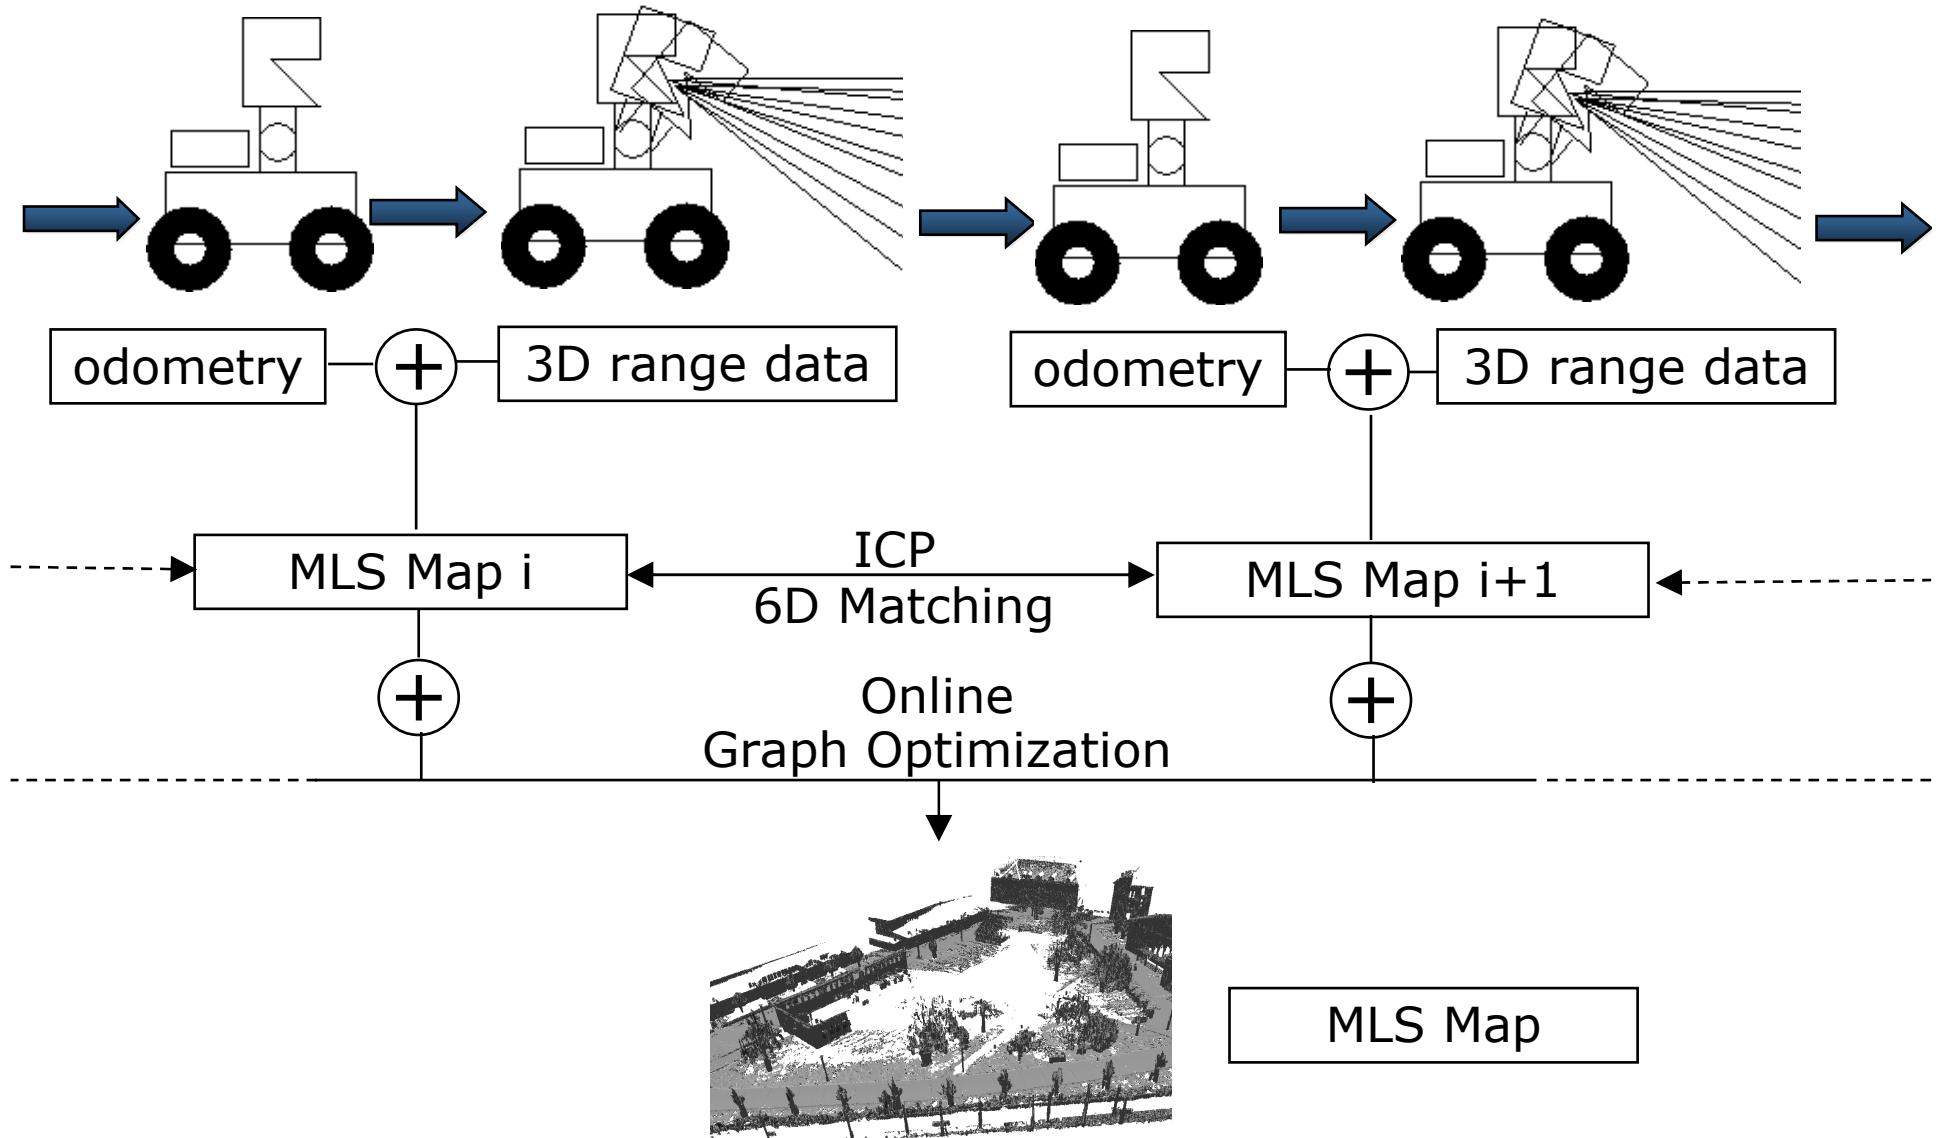
Learning 3D Maps with Laser Data

- Laser range data on a pan-tilt-unit
- Robot that provides odometry

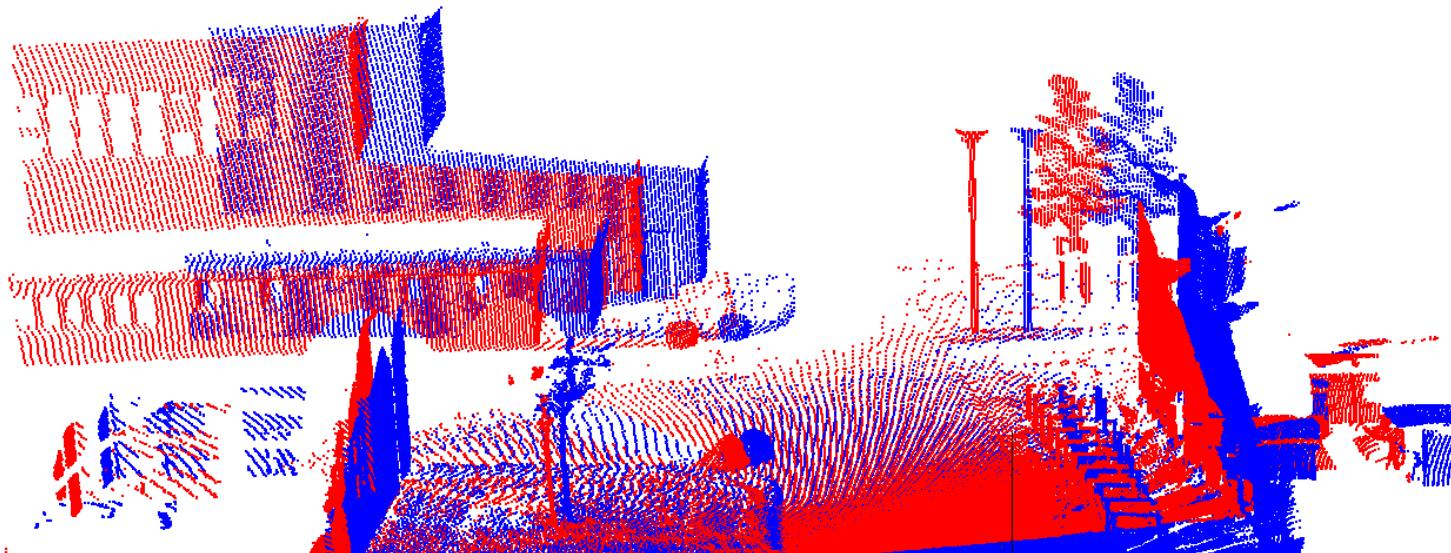


[Kuemmerle et al.]

Incremental 6D SLAM



3D SLAM: Aligning Consecutive Maps



3D SLAM: Aligning Consecutive Maps

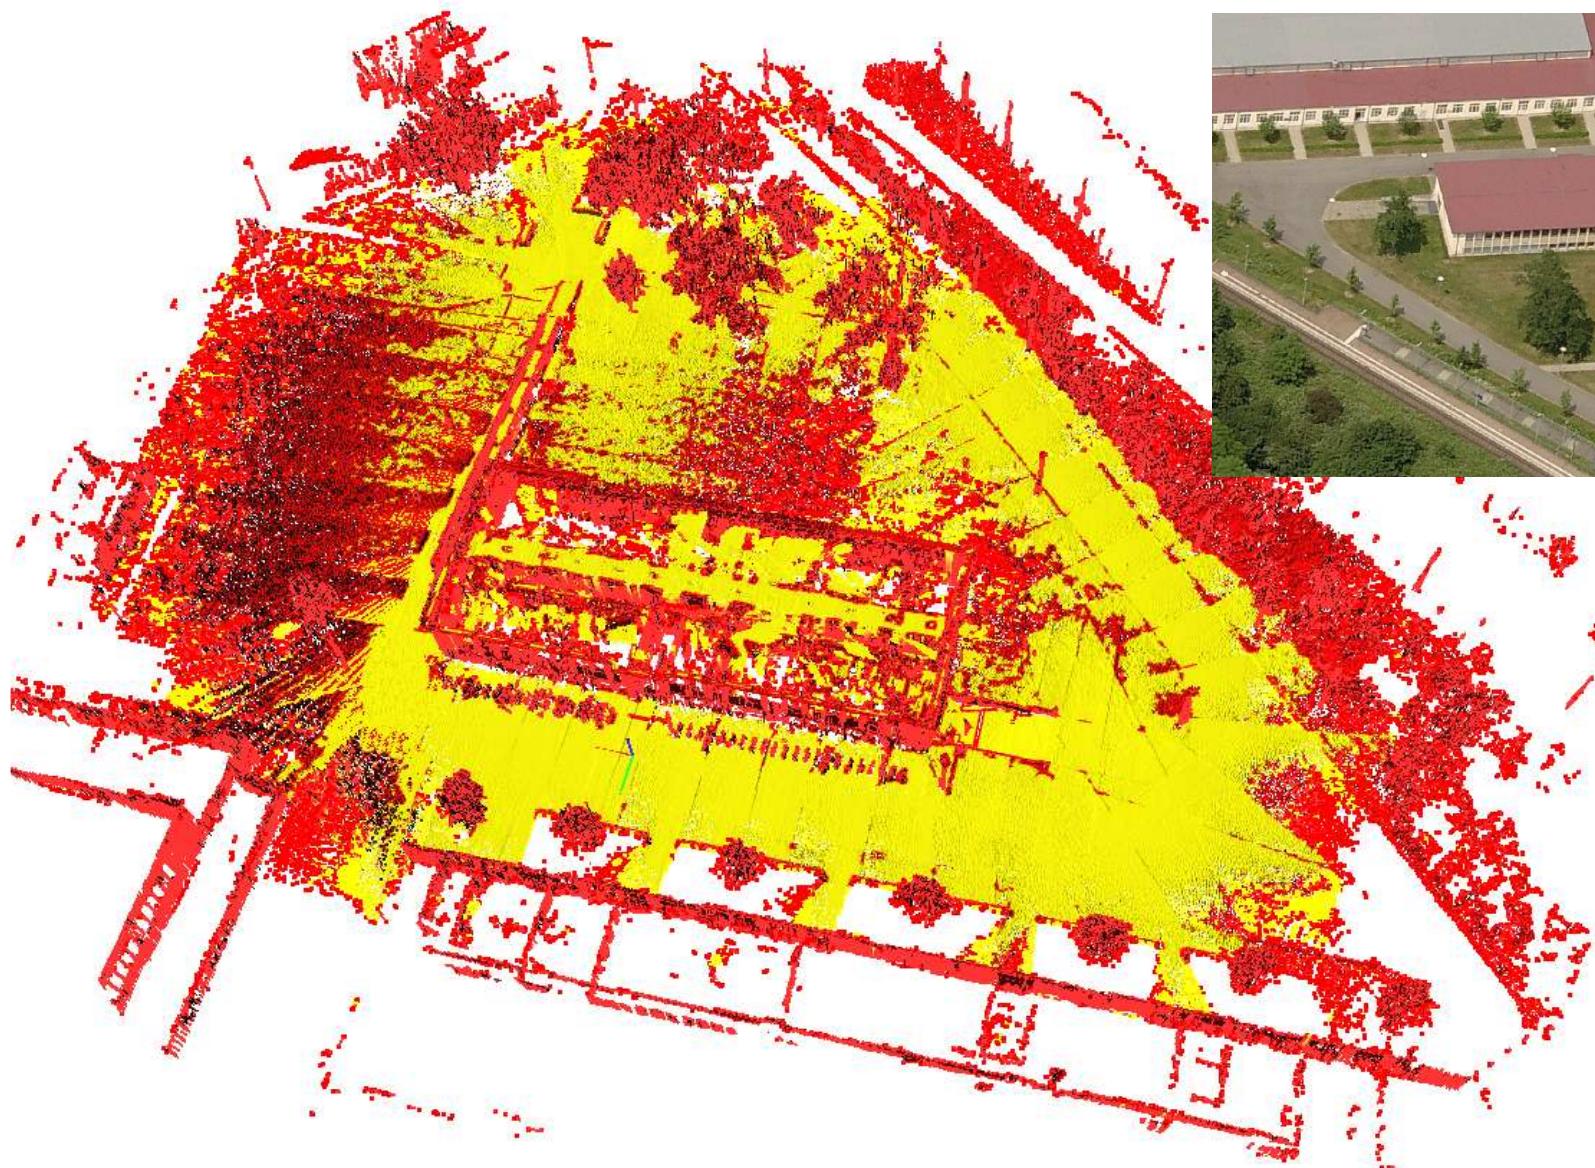
- Given that \mathbf{u}_{i_c} and \mathbf{u}'_{j_c} are corresponding points.
- Try to find the parameters R and t which minimize the sum of the squared error $e(R, t)$

$$e(R, t) = \underbrace{\sum_{c=1}^{C_1} d_v(\mathbf{u}_{i_c}, \mathbf{u}'_{j_c})}_{\text{vertical objects}} + \underbrace{\sum_{c=1}^{C_2} d(\mathbf{v}_{i_c}, \mathbf{v}'_{j_c})}_{\text{traversable}} + \underbrace{\sum_{c=1}^{C_3} d(\mathbf{w}_{i_c}, \mathbf{w}'_{j_c})}_{\text{non-traversable}}$$

Aerial Image of the Scene

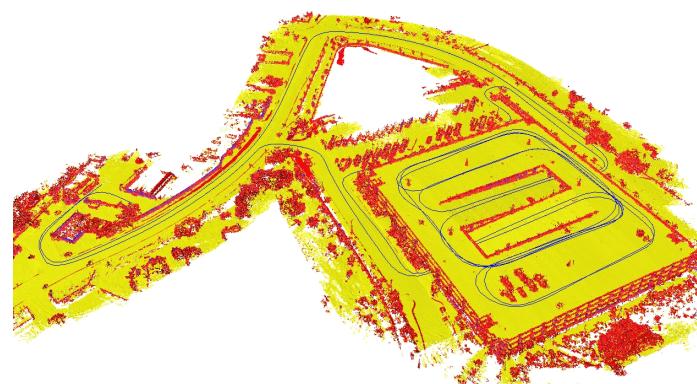


Online Estimated MLS Map



Learning 3D Maps with Laser Data from a Car

- Car robot with a Velodyne 3D laser range scanner
- Use resulting map for autonomous driving



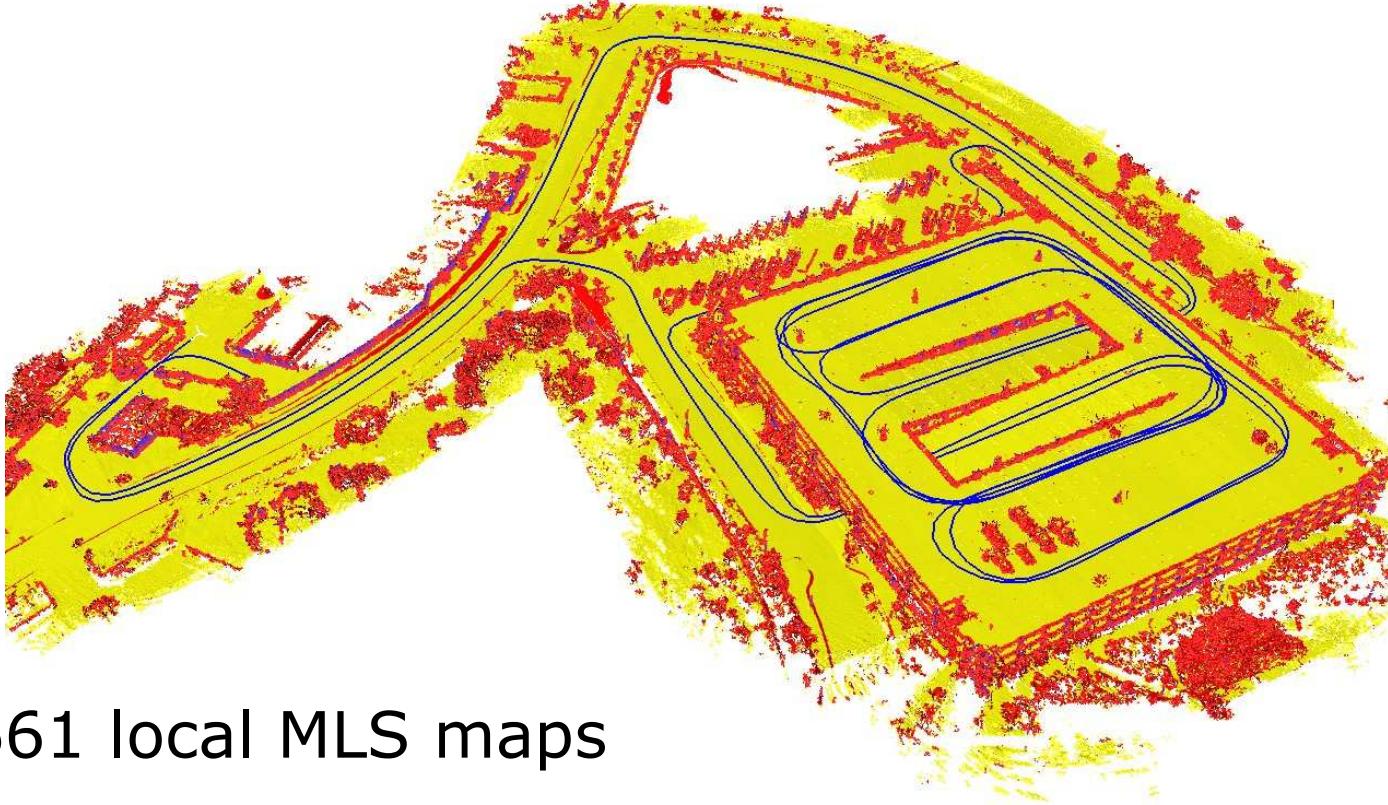
[Kuemmerle et al.]

Learning Large 3D Maps for Driving



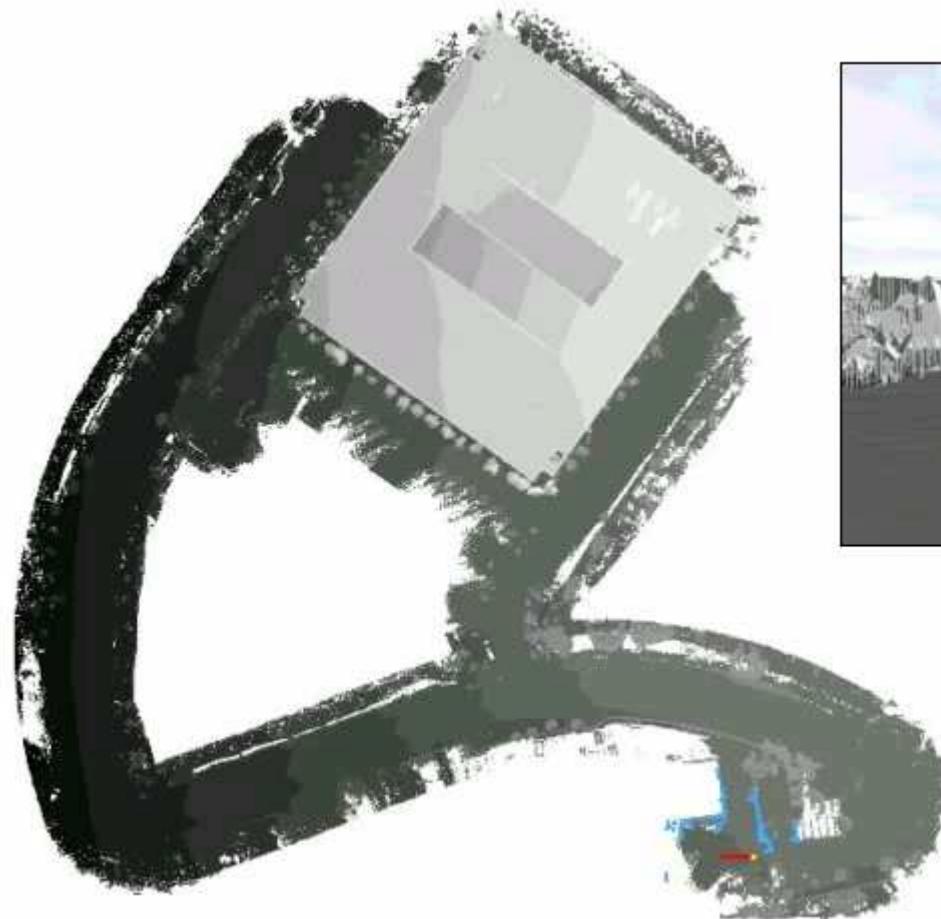
- Parking lot at Stanford University
- Nested loops, trajectory of ~7,000m

Learning Large Scale MLS Maps with Multiple Nested Loops



- 1661 local MLS maps
- nested loops with a total length of ~7,000m
- cell size of 20cm x 20cm.
- requires 118 MB of memory for storage

Application: Car Localization



Learning Maps for Aerial Vehicles

- Learn a map using

- a flying vehicle



- Camera(s)



- An inertial measurement unit



[Steder et al., 08]

Examples of Camera Images



Concrete/Stone



Intersections
concrete/grass



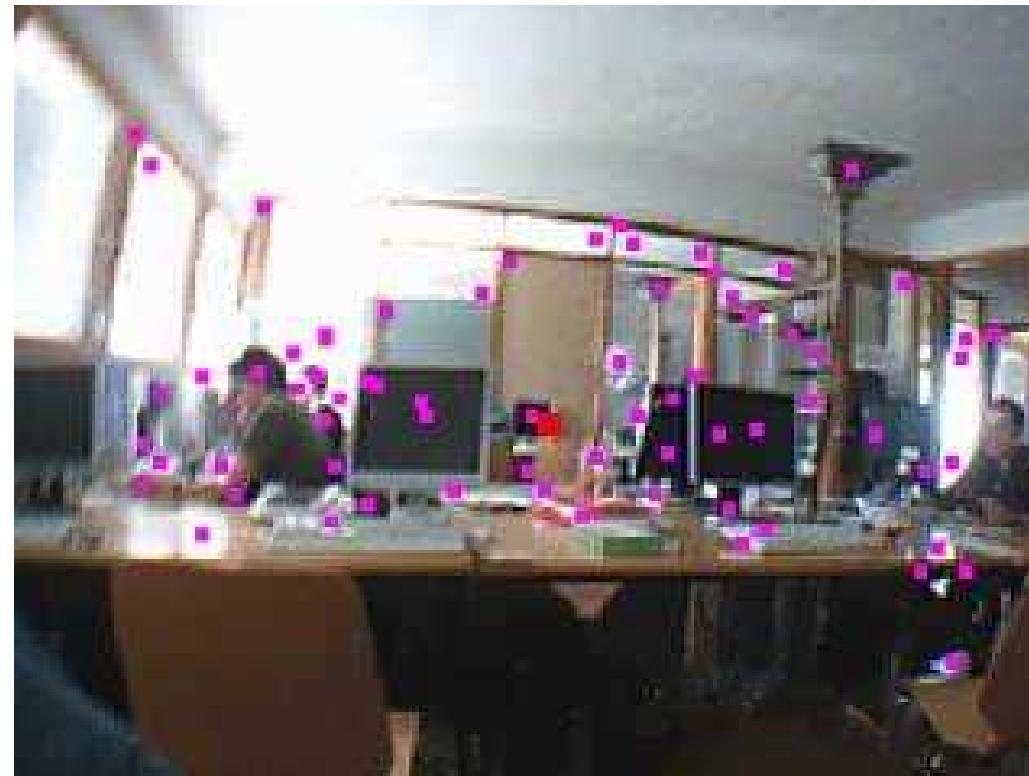
Grass



Wooden floor
(indoor)

Feature Extraction: SURF

- Provide a **description vector** and an orientation
- Rotation and scale invariance



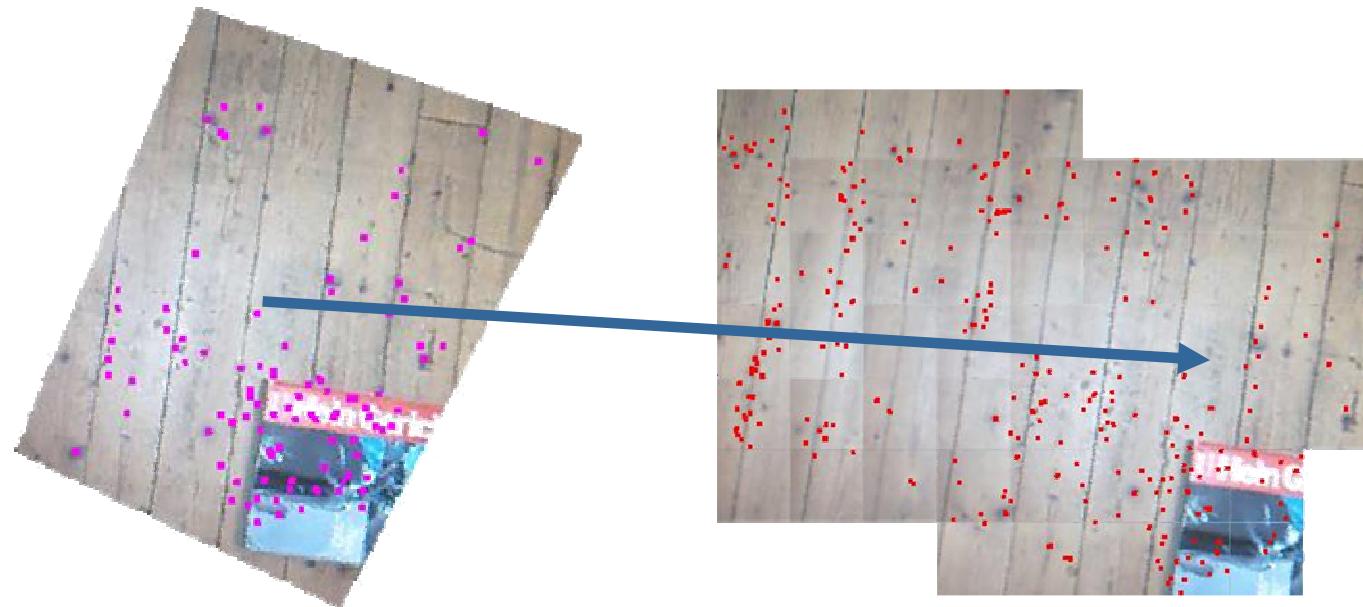
[Bay et al., 06]

Determining the Camera Pose

Wanted: $x, y, z, \varphi, \theta, \Psi$ (roll, pitch, yaw)

- IMU determines roll and pitch accurately enough
 - x, y, z and the heading (yaw) have to be calculated based on the camera images
- 3D position of **two** image features is sufficient to determine the camera pose.

Building a Map



Features in image

Features in map

Match features to extract the position

Camera Pose Estimation

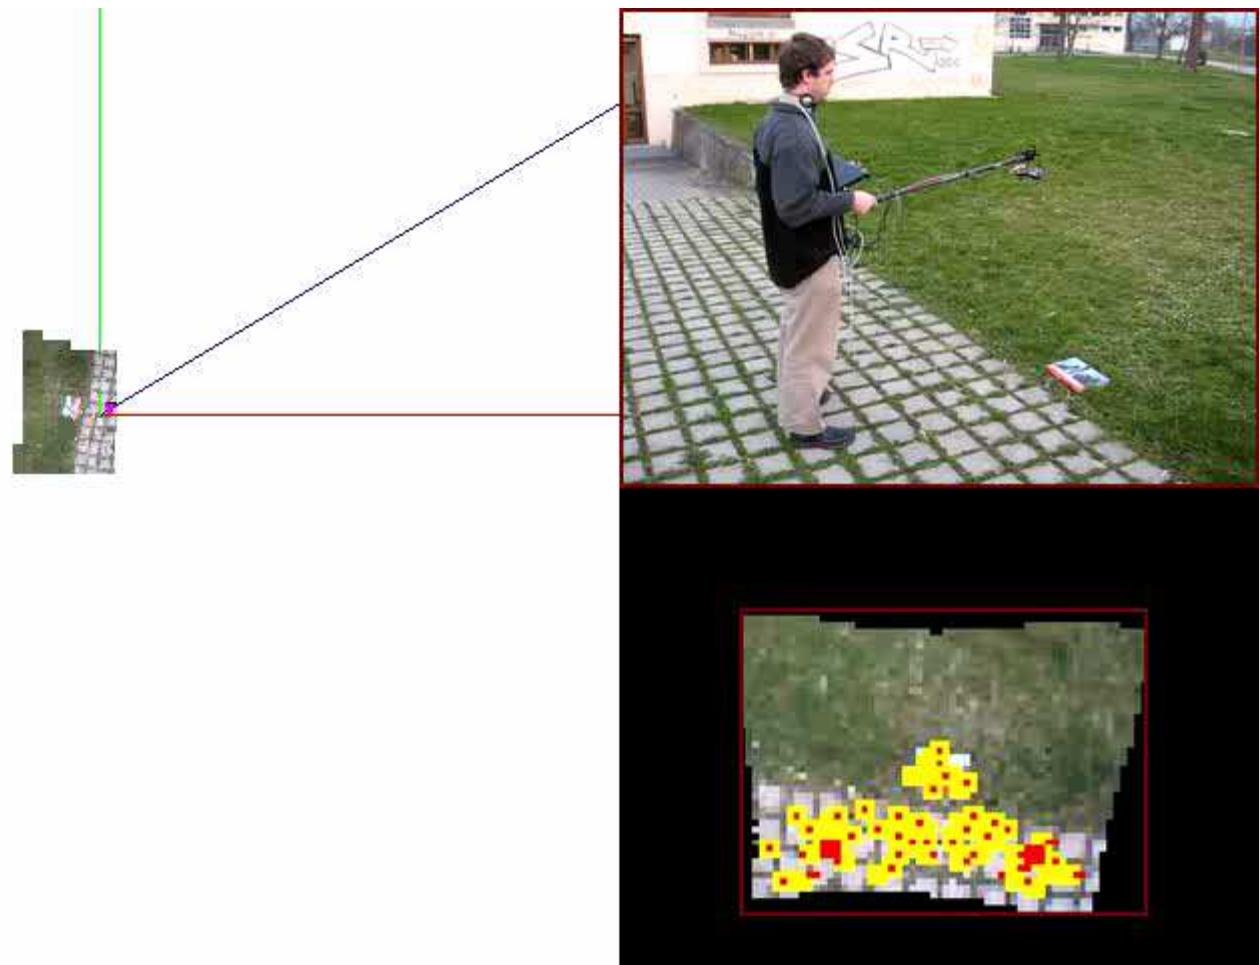
1. Find possible matches.
2. Use the descriptor distance to order matches.
 - Use two matches to calculate the camera position. Start with the best.
 - Re-project all features accordingly to get a quality value about this pose.
 - Repeat until satisfactory pose is found.
3. Update map.

Similar to RANSAC

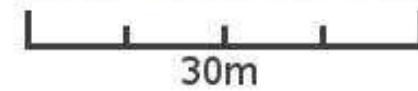
Finding Edges in the Graph

- **Visual odometry:** Compare with temporal close features (e.g., the last 5 frames).
- **Localization:** Compare with features from the map in a given region around the odometry position (local search).
- **Loop closing:** One or more reference features are compared to all map features. Hits result in a localization attempt in that area.

Outdoor Experiment

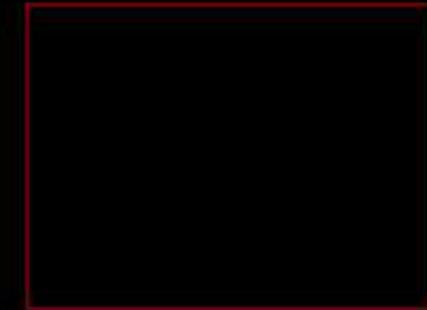


Resulting Trajectory

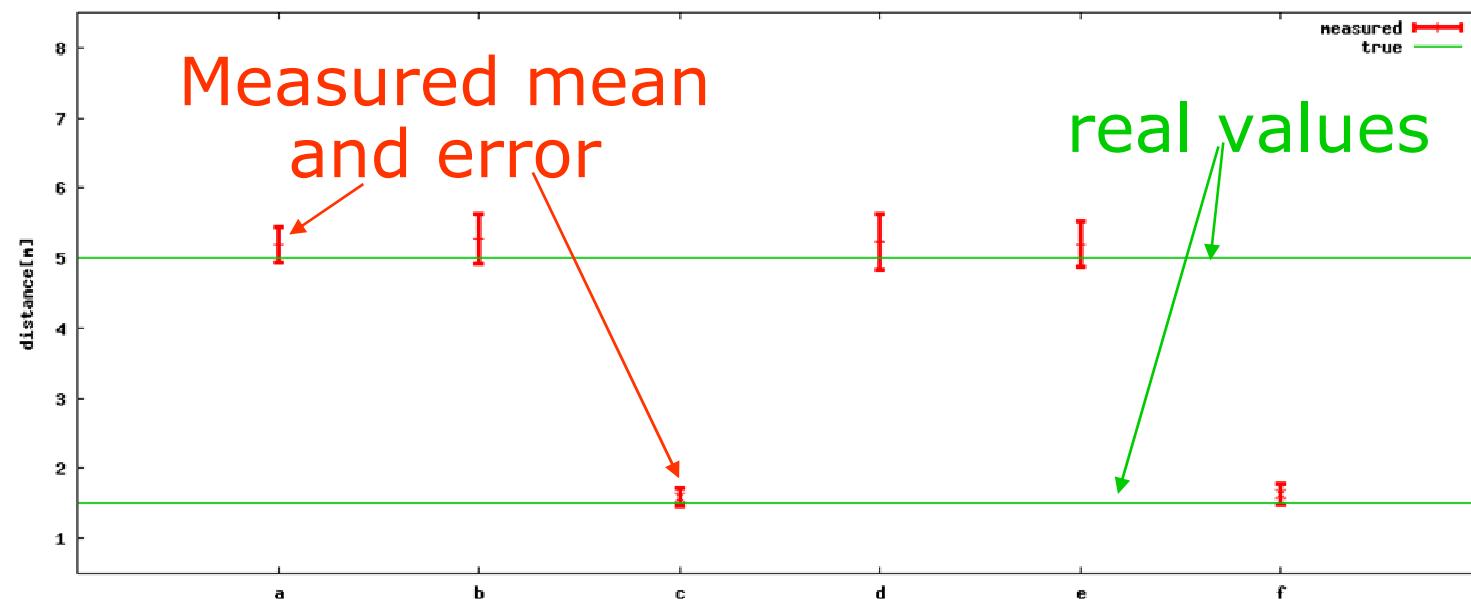
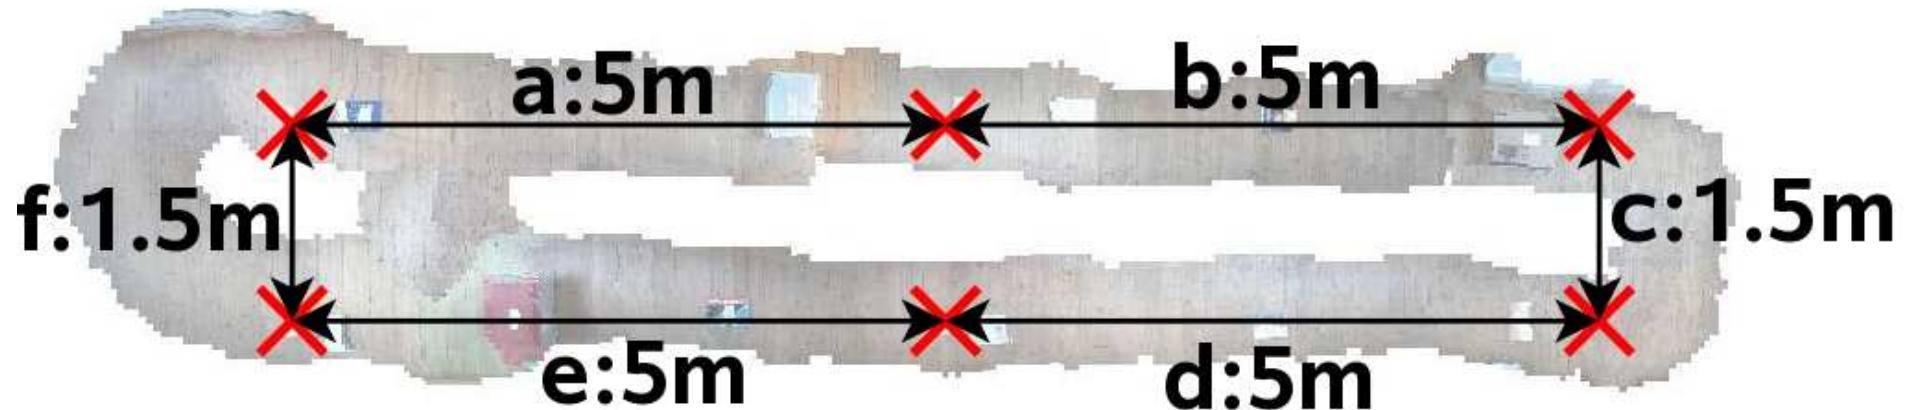


- Length (Google Earth): 188m
- Determined length: 208m

Indoor Experiment



Ground Truth



Result of 10 Loops

Real trajectory length:	23.00m
Average length in map:	24.11m
Standard deviation:	1.32m
Average error:	5.2%
Error in object heights:	16.9%

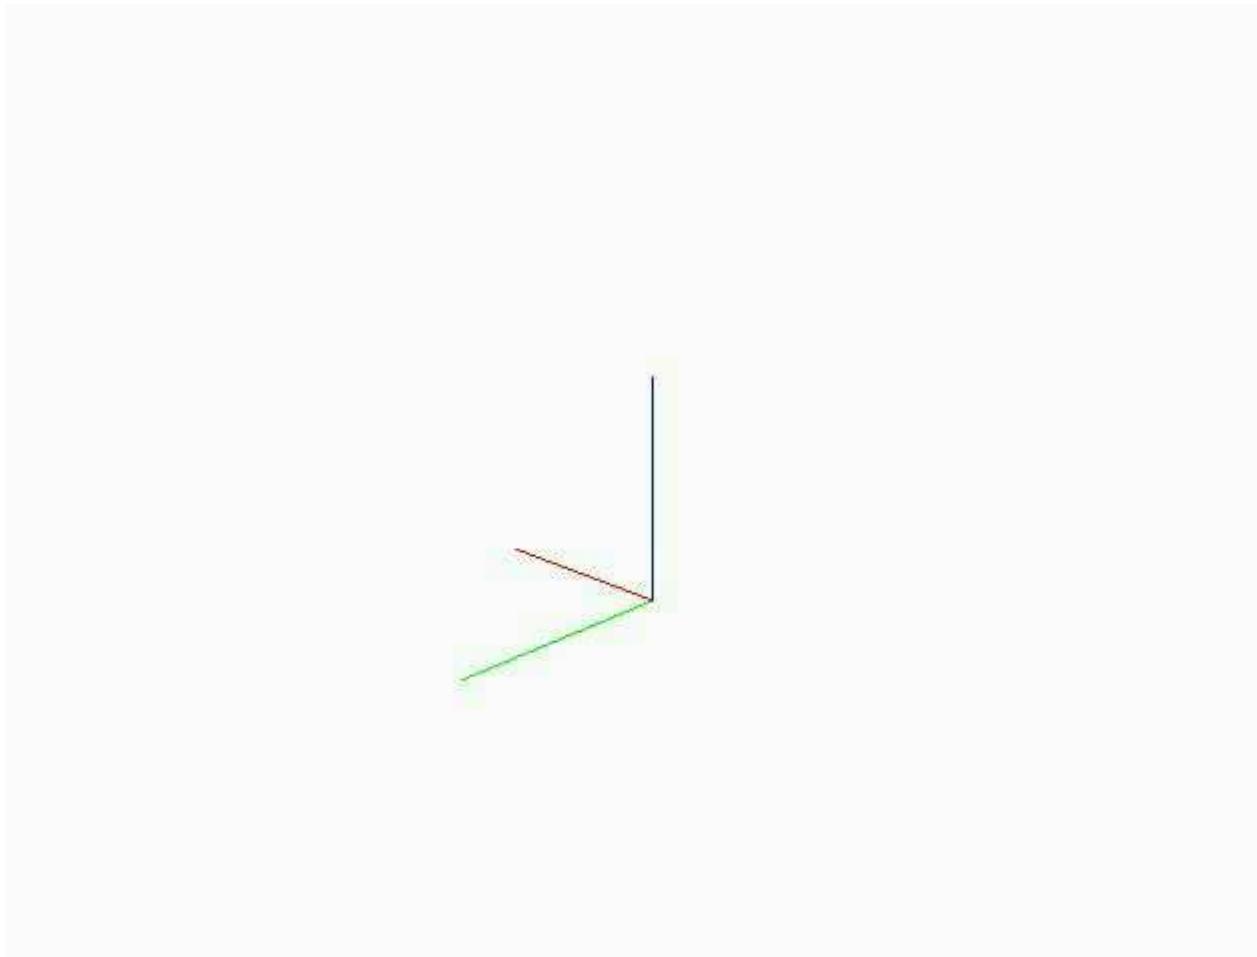
Sources of errors:

- Sensor noise
- Inaccuracy in camera calibration
- Wrong feature matches

Experiment with a Blimp



Experiment with a Helicopter



Conclusions

- Highly efficient approach for optimizing 2D and 3D pose graphs
- Orders of magnitude faster than standard nonlinear optimization approaches
- Covariance estimates can be recovered by means of belief propagation with covariance intersection
- Application examples (standard wheeled robots, autonomous cars, flying vehicles)